

# **From the Edge to the Cloud: Exploring AI Inference Across the Computing Continuum**

(yes, including Generative AI)



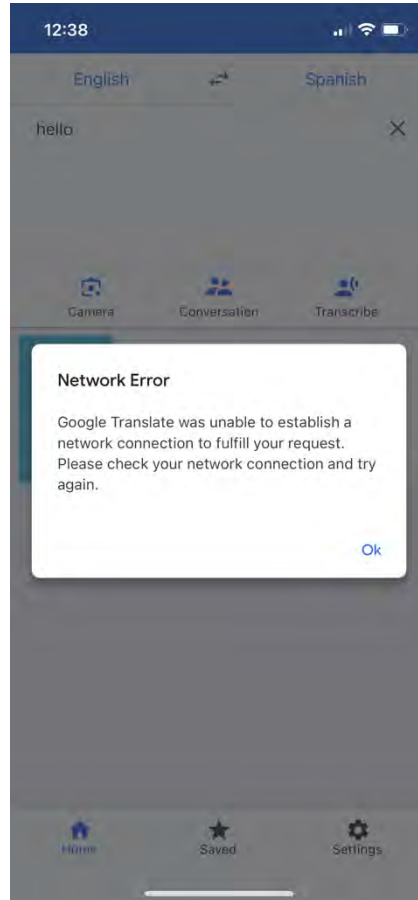
Roberto Morabito  
Assistant Professor @ EURECOM  
<https://www.linkedin.com/in/robertomorabito>

# From the Edge to the Cloud: Exploring AI Inference Across the Computing Continuum

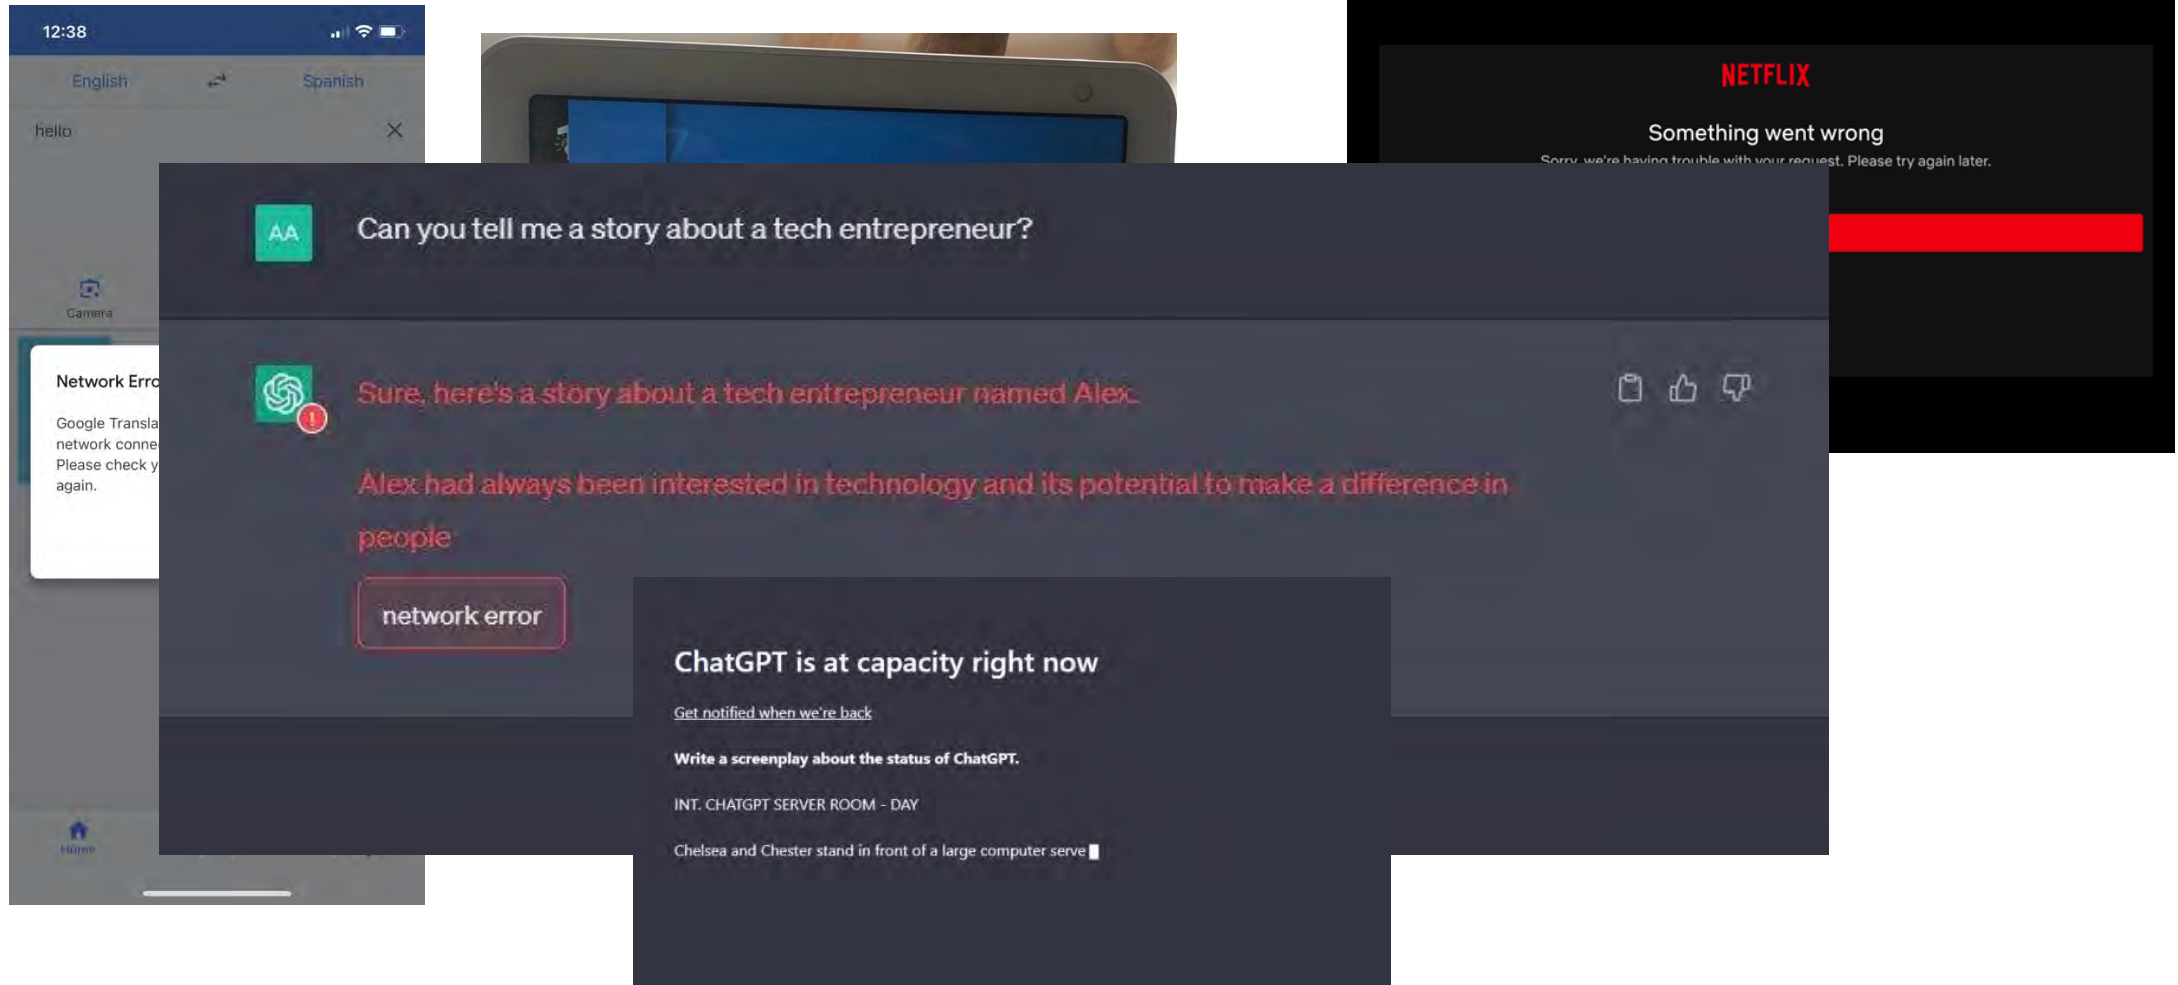
(yes, including Generative AI)

**Let's start from the edge**


# Cloud Computing Is Great **But...**



# Cloud Computing Is Great **But...**



# Edge Computing



- Idea is to **push applications, data and computing power to the edge of the Internet**, near mobile devices, sensors, and end users

## Main Drivers

### Latency

- Data processing close to where it originates avoids round-trip time to the cloud

### Bandwidth

- Optimization of communication to and from the cloud


### Privacy / Security

- Sensitive data stays local


### Connectivity

- Continued processing (in some cases) despite lack of connectivity to the cloud

# Edge AI




# Edge AI Chips MARKET



“Bringing AI to the device: Edge AI chips come into their own “ Source: <https://bit.ly/3r31Jv>

# Edge AI Chips – AI Acceleration



Source: <https://www.thinkautonomous.ai/blog/vision-processing-units-vpus/>



# AI Inference At The Edge 2019 Vs. 2024

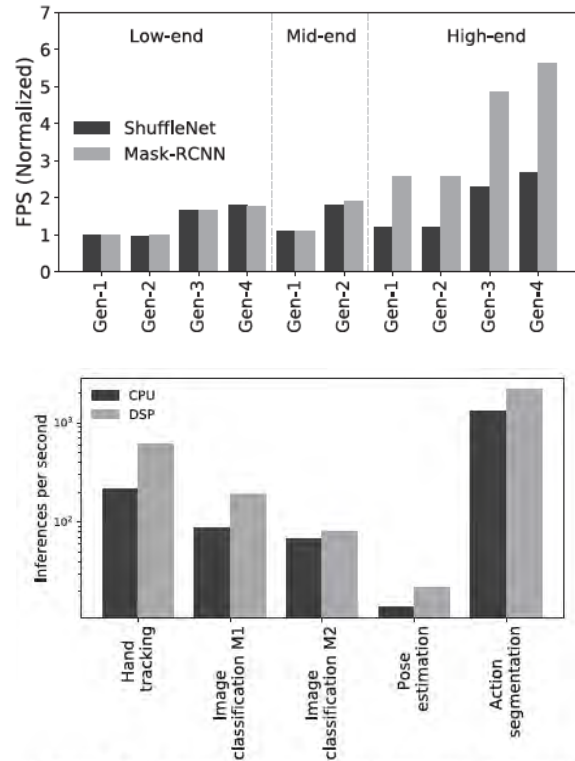


Figure 8: Inference time performance comparison between CPU and DSP.

(Source: Wu, C.J., Brooks, D., Chen, K., Chen, D., Choudhury, S., Dukhan, M., Hazelwood, K., Isaac, E., Jia, Y., Jia, B. and Leyvand, T., 2019, February. Machine learning at facebook: Understanding inference at the edge. In 2019 IEEE international symposium on high performance computer architecture (HPCA) (pp. 331-344). IEEE.)



## Samsung bets heavily on AI tricks to boost Galaxy S24 appeal

South Korean firm will hope generative AI text, voice, image and video tools can help it regain top spot in phone market



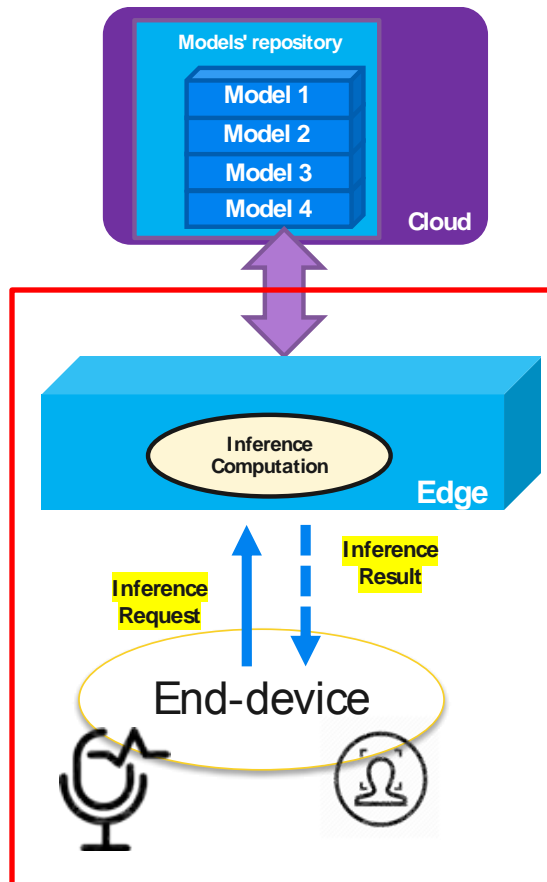
## Samsung unveils Galaxy S24 series with AI-powered Camera, Translation, and Editing tools

Samsung says users can decide how much data they want to use for AI features, and opt out of online processing if they want to




(Source: <https://www.theguardian.com/technology/2024/jan/18/samsung-bets-heavily-on-ai-tricks-to-boost-galaxy-s24-appeal> and <https://mobilesyrup.com/2024/01/17/samsung-s24-series-ai-features/>)

# AI Inference In Distributed Edge Computing Systems

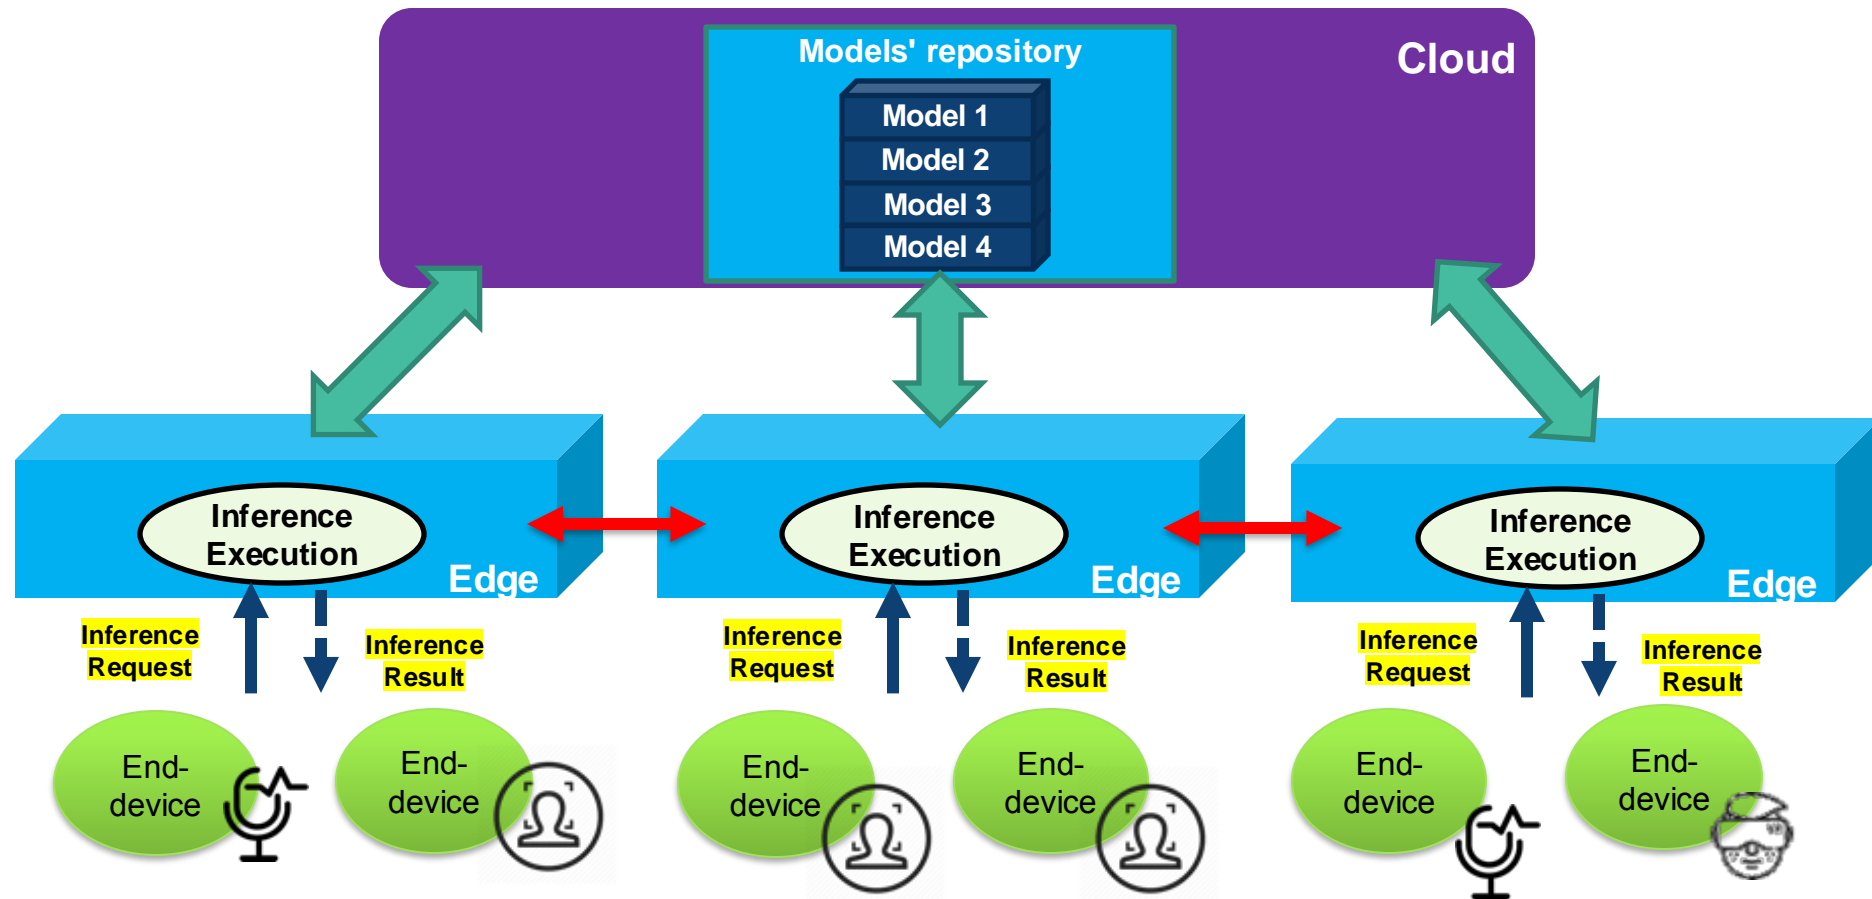


From the perspective of distributed edge AI systems, related studies had primarily focused on theoretical models and simple scenarios involving interactions between a single device or edge and the cloud.

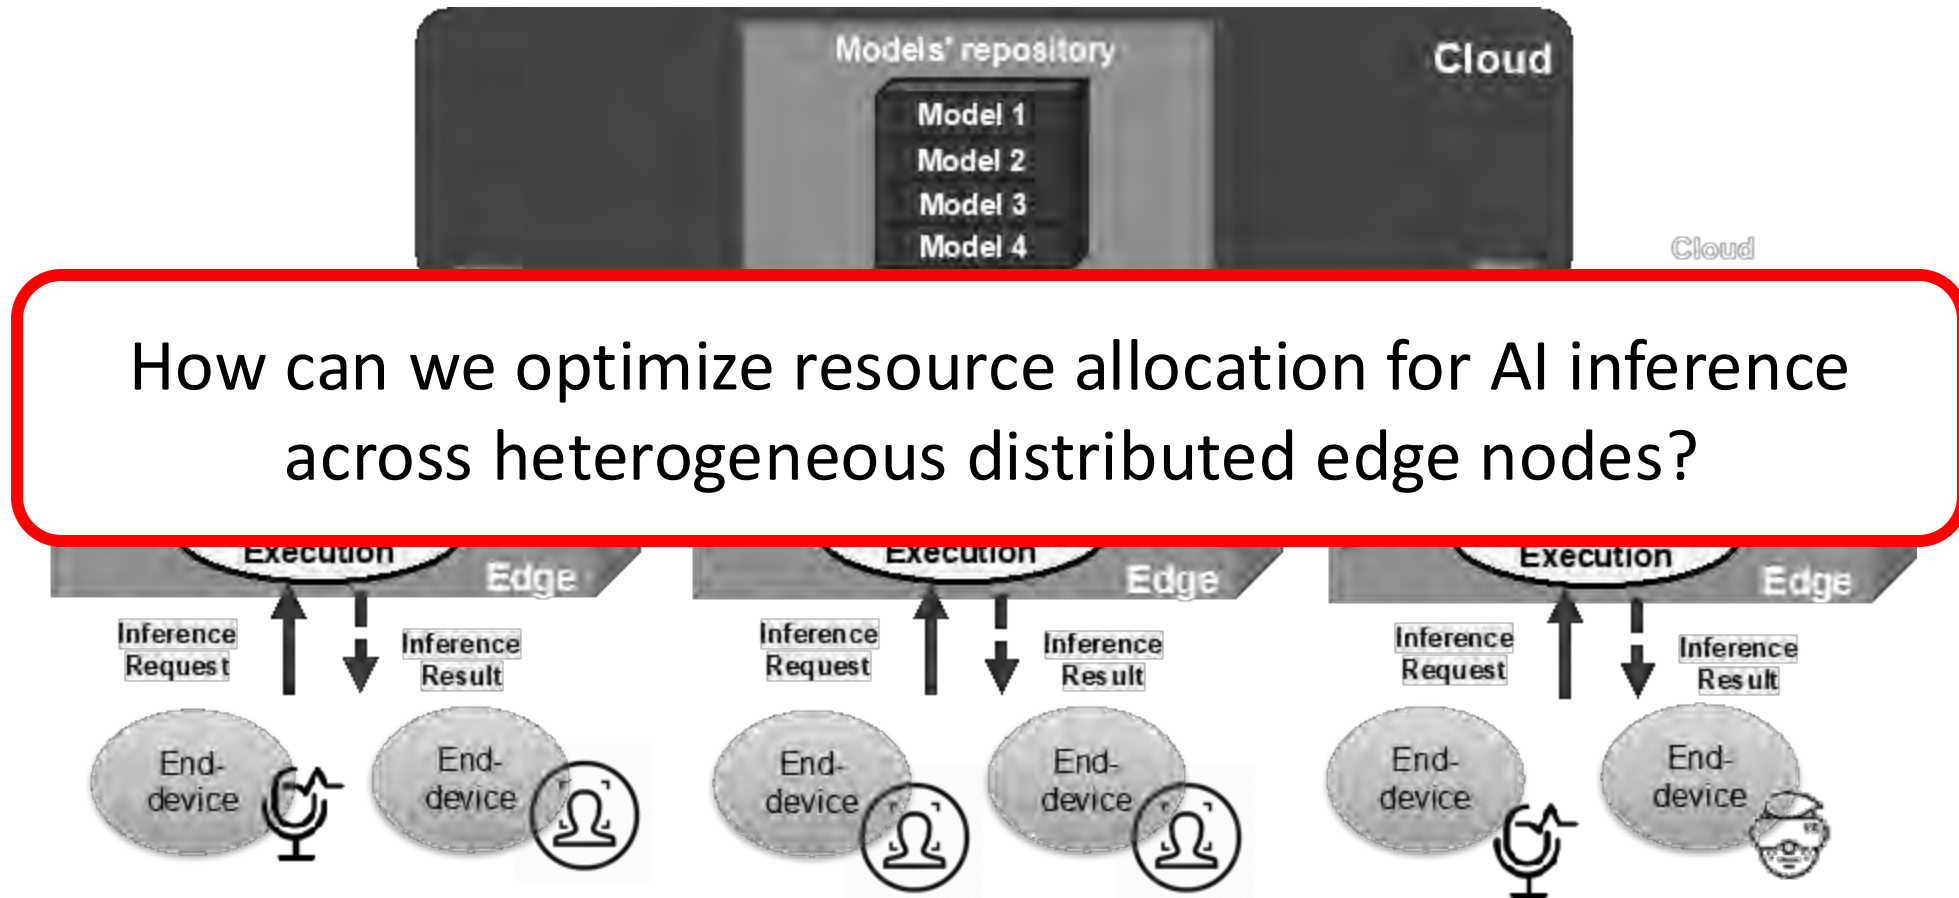
- Not many scenarios where **multiple edge nodes** are involved
- **Hardware heterogeneity** and **networking aspects** are often not considered



# AI Inference In Distributed Edge Computing Systems



# AI Inference In Distributed Edge Computing Systems



# Challenges And Requirements: Latency

## Latency Components

- Composed of communication latency (data exchange) and computing latency (model training/inference execution).

## Latency Significance

- Crucial for Inference: Requires **near real-time execution for prompt responses**.

Examples:

- Voice assistants need predictions within **200ms**.
- Tactile Internet and autonomous driving operations demand below **10ms** latency.



Source: Campolo, C., Iera, A. and Molinaro, A., 2023. Network for Distributed Intelligence: a Survey and Future Perspectives. IEEE Access.

# Challenges And Requirements: Reliability

## Reliability Components

- The ability of the network to consistently perform its intended function accurately and dependably.
- Key Aspects: Includes error rates, **uptime**, and **fault tolerance**.


## Reliability Significance

- Crucial for Consistent AI Performance: Ensures that **AI systems can function correctly** and deliver accurate results over time, **regardless of network conditions**.
- Impact on AI Applications: **High reliability** is essential for **mission-critical AI applications where errors or downtime** can have severe consequences.




# Challenges And Requirements: Practical Development

Realistic and more complex Edge AI/IoT deployment scenarios demand additional **requirements** to be fulfilled.

- ***Plug and Play interoperability*** among edge devices embedding different AI accelerators (e.g., **GPU, VPU, TPU**).
- **Agnostic AI inference services discovery** and **provisioning**.
- Combined **computing-** and **networking-** aware **orchestration mechanisms**, suitable for satisfying the Quality of Service (QoS) requirements of AI-enabled applications



# (Resource Constrained) Heterogeneous Edge AI Nodes

Feature	<b>Coral Dev Board</b> 	<b>Jetson Nano</b> 	<b>Up Squared AI Edge X</b> 
CPU Chipset	NXP i.MX 8M SoC (quad Cortex-A53, Cortex-M4F)	Quad-core ARM Cortex-A57 MPCore processor	Intel® Apollo Lake SoC ATOM x7-E3950
AI Accelerator Chipset	Google Edge <b>TPU</b> coprocessor: 4 TOPS (int8)	<b>GPU</b> NVIDIA Maxwell architecture with 128 NVIDIA CUDA® cores	Movidius Myriad X <b>VPU</b> integrated
Memory RAM	1GB LPDDR4	4GB LPDDR4	8GB LPDDR4
Storage	8GB eMMC MicroSD card slot	MicroSD card slot	64GB eMMC
Connectivity	Ethernet (1 x GbLAN) WiFi Bluetooth	Ethernet (1 x GbLAN) WiFi	WiFi 802.11 AC 2T2R + Bluetooth 4.2 (BLE)+ *LTE + Gigabit Ethernet




# Hardware And Software Heterogeneity Major Implications

Impossible full '*out of the box*'  
devices' interoperability

AI Inference Latency may vary  
from board to board

Lack of seamless AI Inference  
provisioning offloading



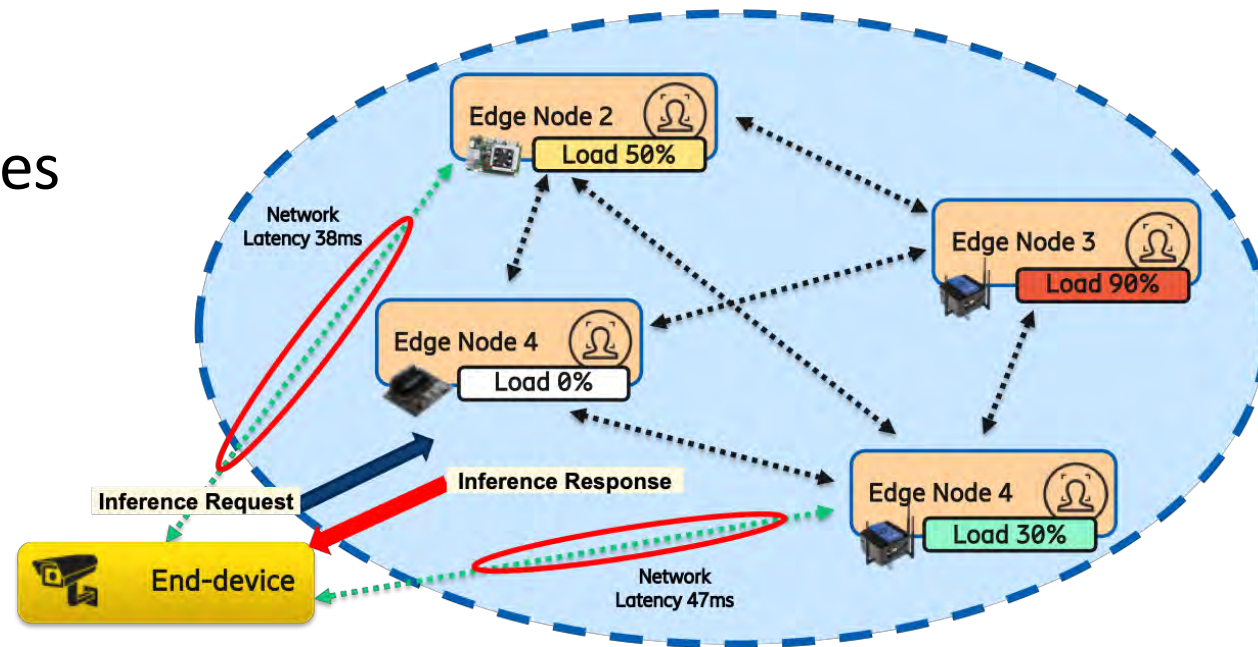
# AI Inference Provisioning In Distributed Edge Systems

## Edge Nodes Resources

- Computation Capabilities

## Network Performance

- Latency
- Bandwidth



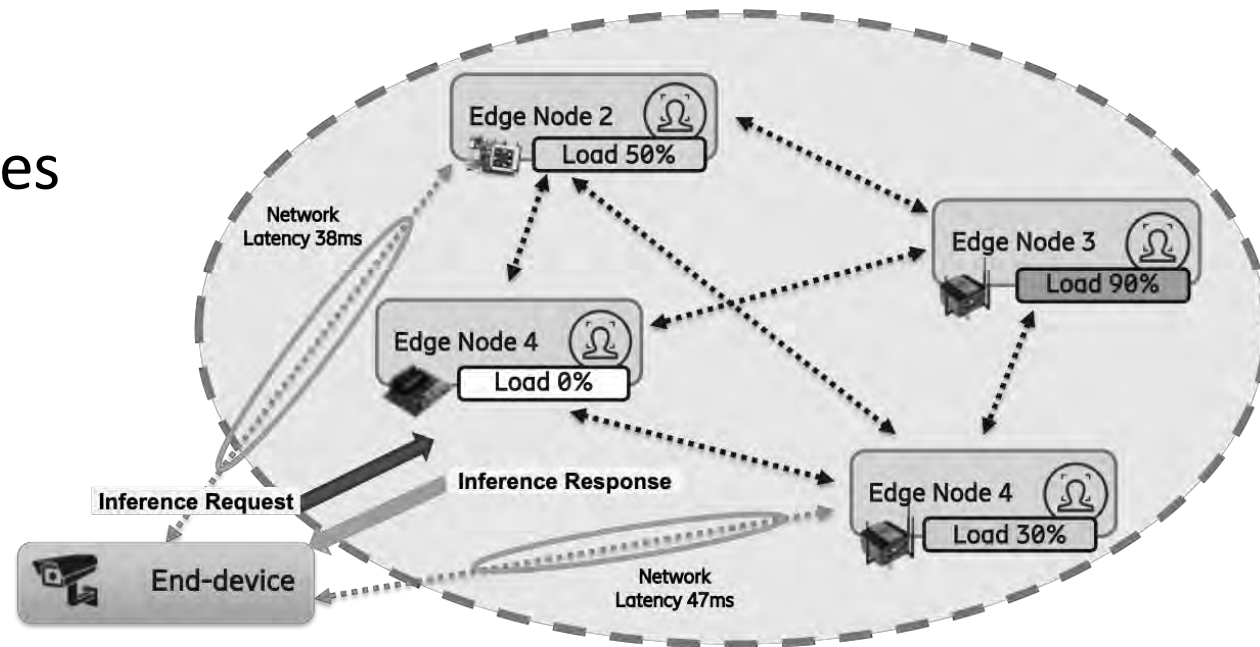
# AI Inference Provisioning In Distributed Edge Systems

## Edge Nodes Resources

- Computation Capabilities

## Network Performance

- Latency
- Bandwidth


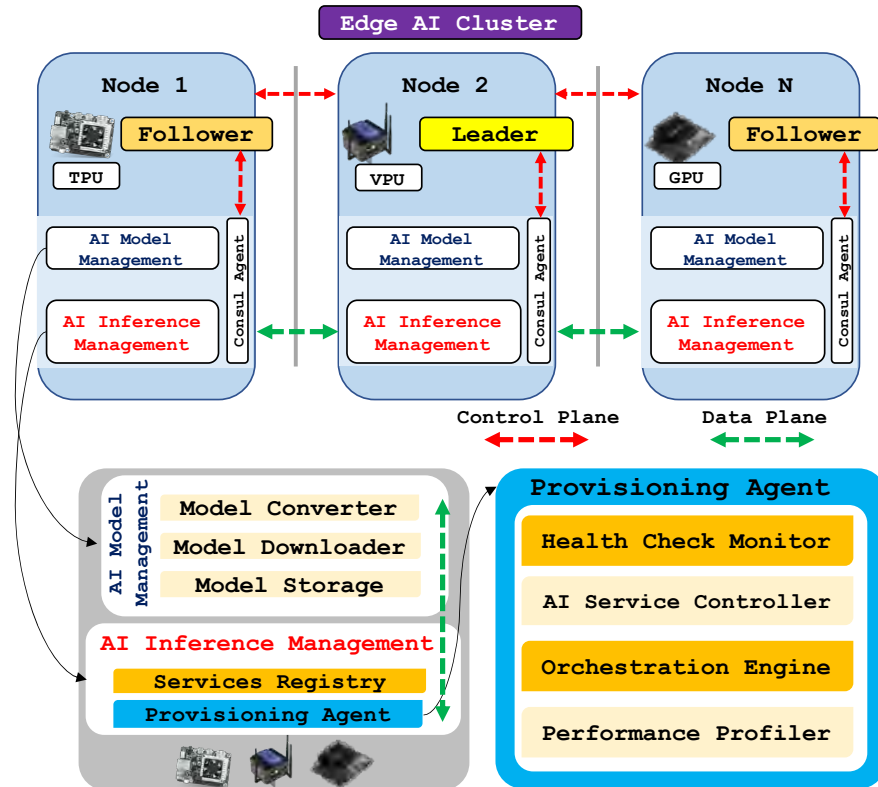


**Which Edge AI node can best provide a specific AI inference service while meeting the requesting device's QoS requirements?**

# AI Inference Provisioning In Distributed Edge Systems

It ensure an abstraction layer that:

- Enables **interoperability** between different AI-enabled devices
- Allows **platform-agnostic service discovery** and provisioning of AI inference services
- Supports **seamless service orchestration** and execution migration capabilities





# Edge AI Testbed

## Edge AI Cluster:

- Intel Movidius Myriad X VPU (UP Squared AI Edge X)
- Google Edge TPU (Coral Dev Board)
- NVIDIA 128-core Maxwell GPU (Jetson Nano)

## End-Devices:

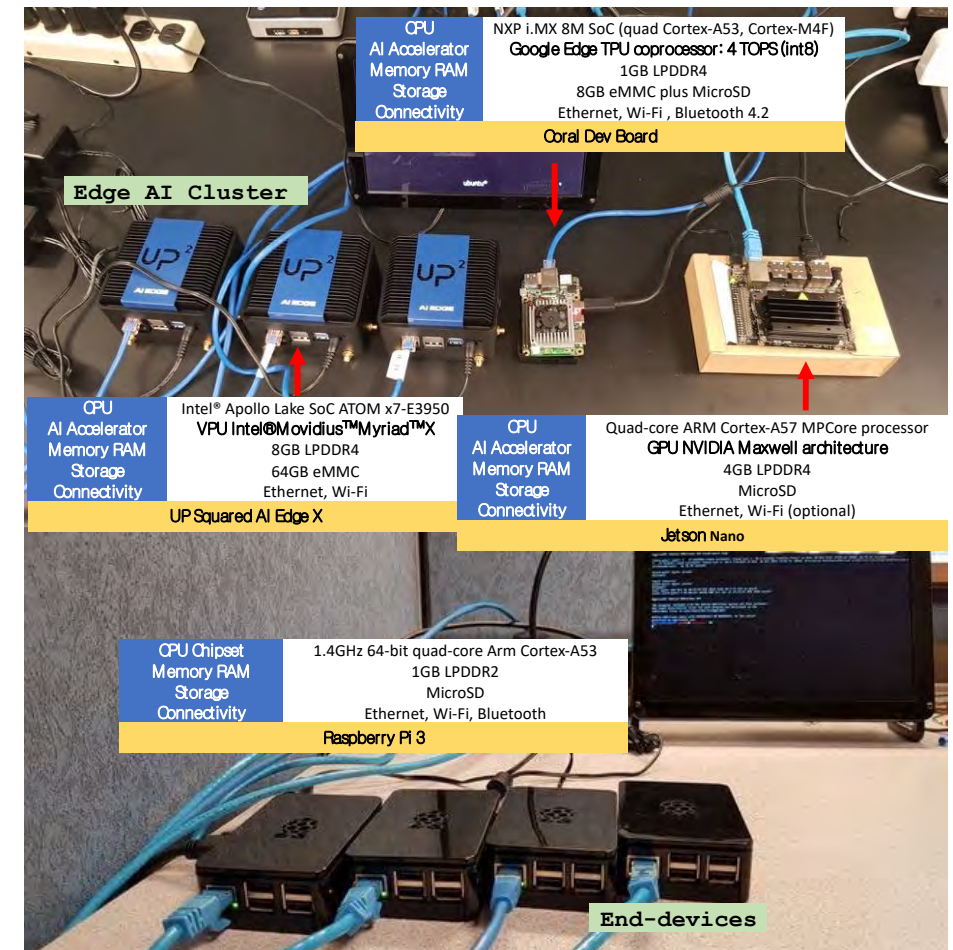
- Raspberry Pi 3 Model B (x4)

## Network Setup:

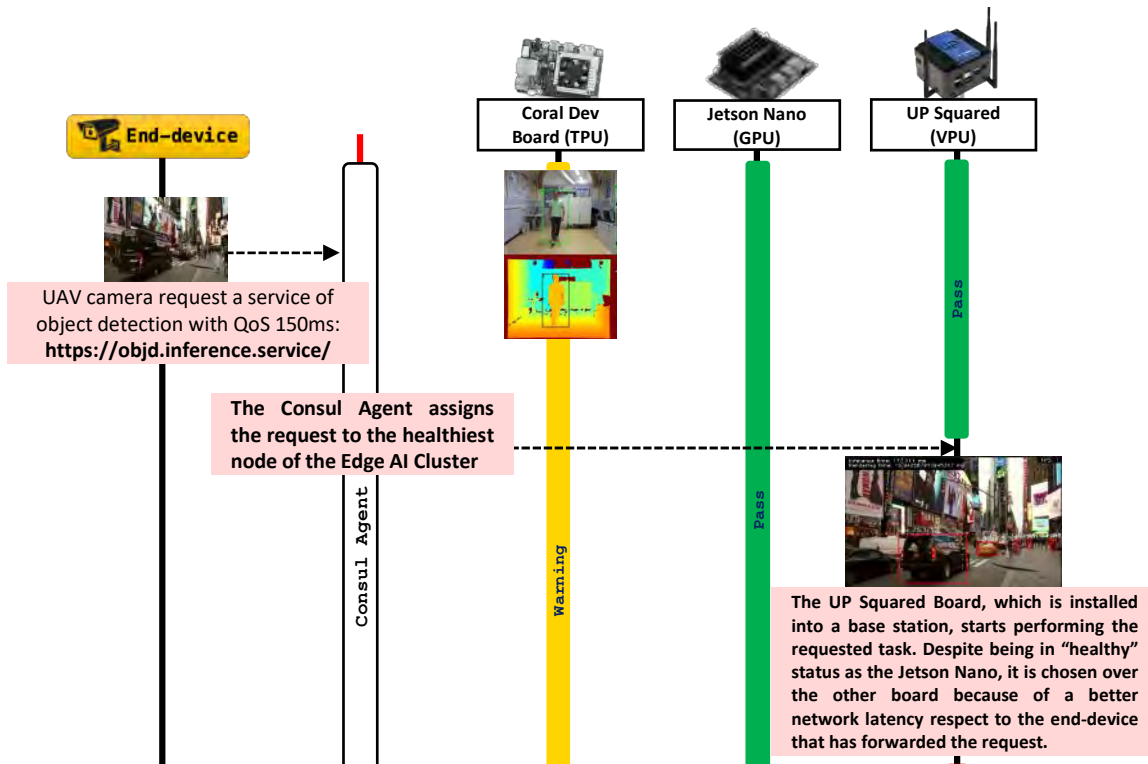
- Controlled wireless network for device communication
- Emulation of realistic edge system deployment

## Latency Measurement & Emulation:

- Analysis based on 5G to edge server latency
- Best-fit distribution: Stable (Shape: 1.6878, Scale: 0.0980)
- Average network latency: 13.405 ms
- Standard deviation: 16.065 ms, showing high variability



# AI Inference Provisioning: Edge AI Node Selection



## Algorithm 1 Select node for AI Inference provisioning and AI Inference execution offload


**Require:** Set of edge nodes (*edgeNodes*), end-device (*endDevice*), NLM (*NetworkLatencyMatrix*), QoS requirements

**Ensure:** Selected healthy edge node (*selectedNode*)

```

1: Preload AI models in edgeNodes
2: Initialize selectedNode to null
3: function ASSIGNNODE(edgeNodes, endDevice, latencyMatrix)
4:   Initialize minLatency to max(NetworkLatencyMatrix)
5:   Initialize healthyNode to null
6:   for each edgeNode in edgeNodes do
7:     if edgeNode is healthy and the latency between edgeNode and endDevice
       is less than minLatency then
8:       Update minLatency with the latency between edgeNode and endDevice
9:       Update healthyNode with the current edgeNode
10:    end if
11:  end for
12:  Assign the AI inference task to edgeNode
13: end function
14: selectedNode ← ASSIGNNODE(edgeNodes, endDevice, NetworkLatencyMatrix)
  
```

# AI Inference Provisioning: Service Provisioning Offloading



## Algorithm 1 Select node for AI Inference provisioning and AI Inference execution offload

```

15: while true do
16:   Continuously monitor the health status and QoS requirements for selectedNode
17:   if selectedNode cannot guarantee required QoS for handled applications then
18:     Temporarily make selectedNode unreachable for new requests
19:     Find another healthy node as described in the ASSIGNNODE function
20:     if a new healthy node is found then
21:       Perform a seamless migration of the affected AI inference task to the
         new healthy node
22:     end if
23:   end if
24: end while
25: return selectedNode

```




*Additional details about the Health Checks definition can be found in the references*



# AI Inference Service Execution Orchestration

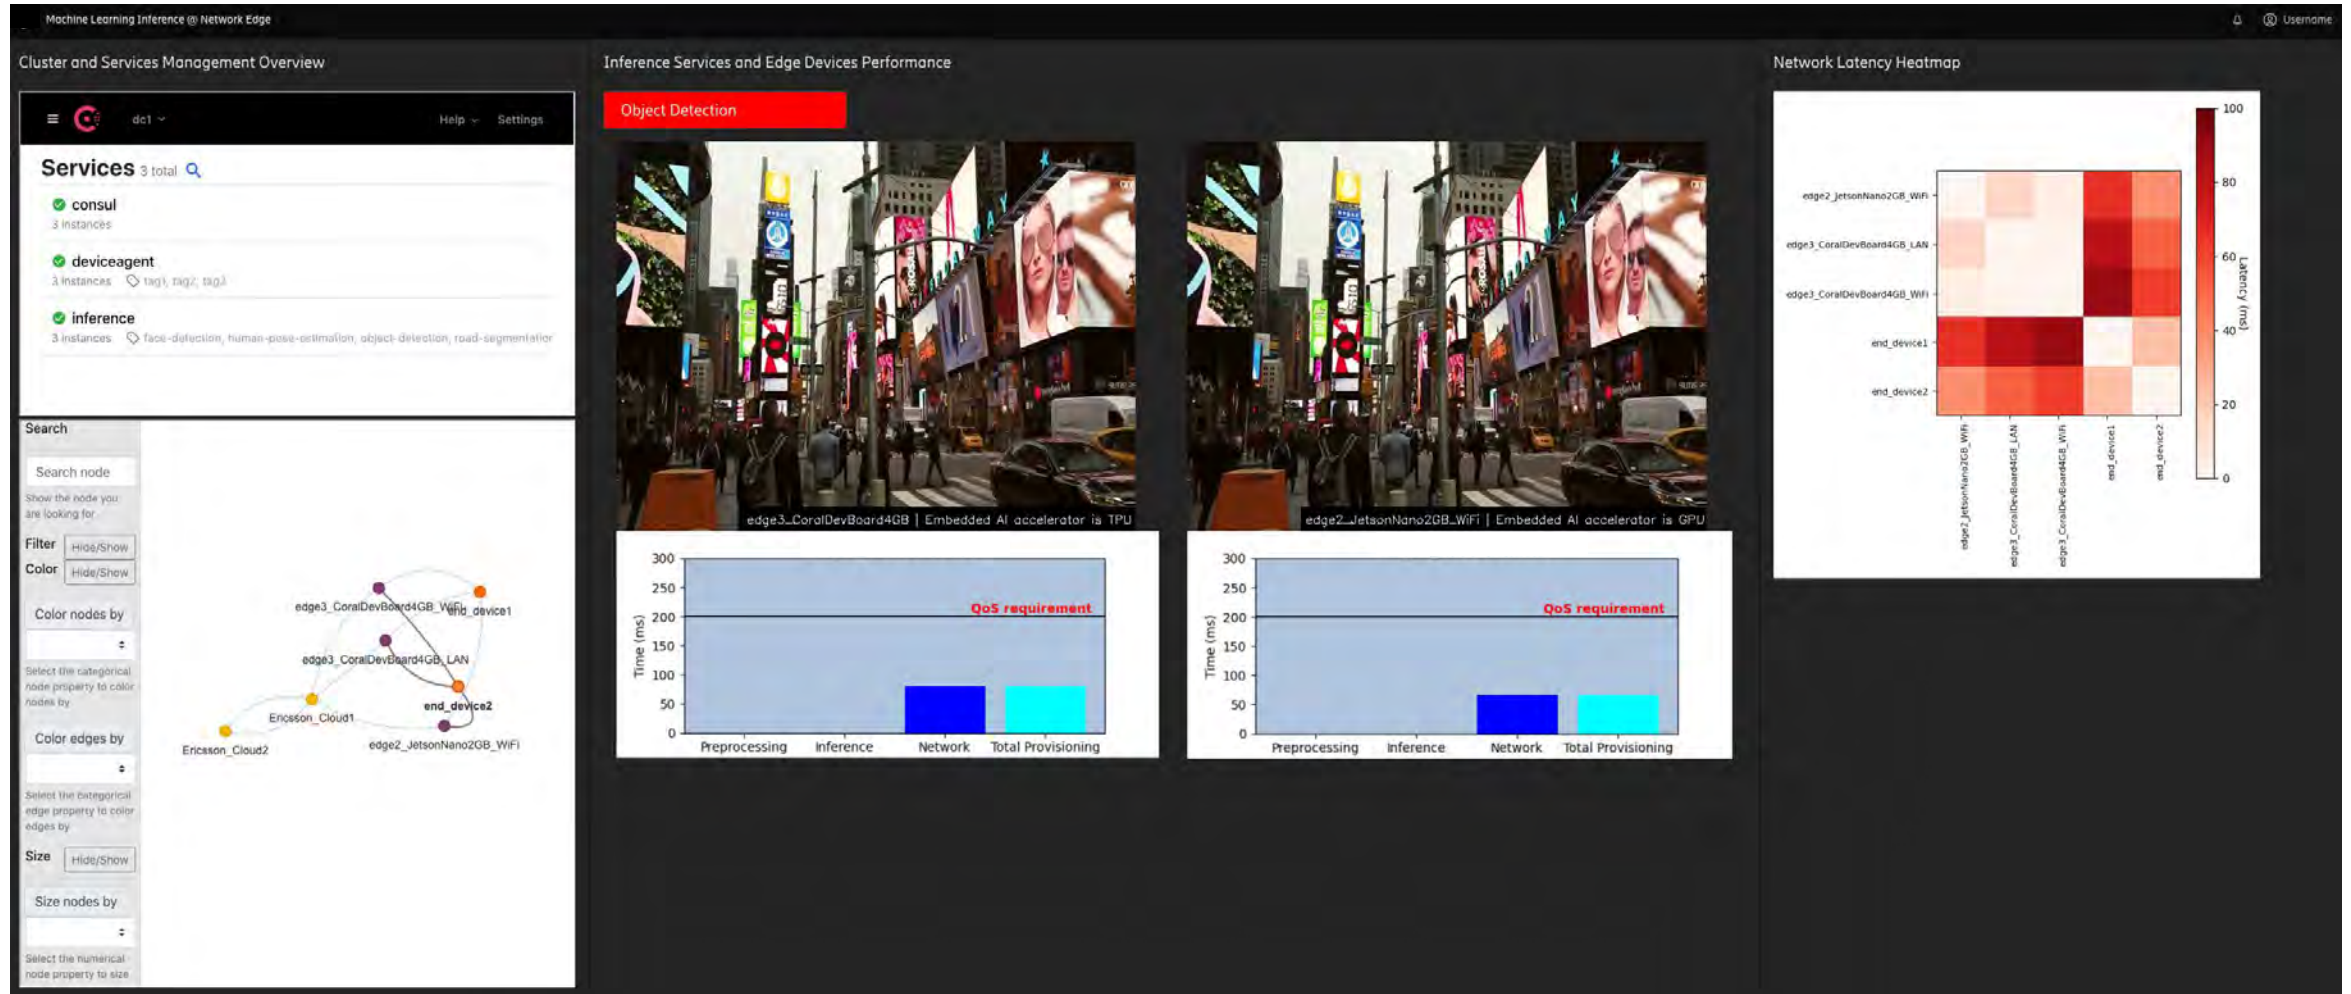
## Object Detection over real-time video streaming

- The application uses an object detection model (MobileNet-SSDv1).
- The ML inference execution is migrated from an edge node to another one as soon as the device gets too overloaded, or network latency increases.







# Demo





# Inference Execution Migration

## Inference Services and Edge Devices Performance

## Object Detection



# Hardware And Software Heterogeneity Impact On The AI Inference Performance



## Insights:

- CPU more affected by frame size increase than AI Accelerator
- Overall system latency inflates by 45-60% with larger frame size
- **AI application parameters significantly affect inference latency, with different devices showing varied responses to the same workload**
- **Device responsiveness to computational demands varies**



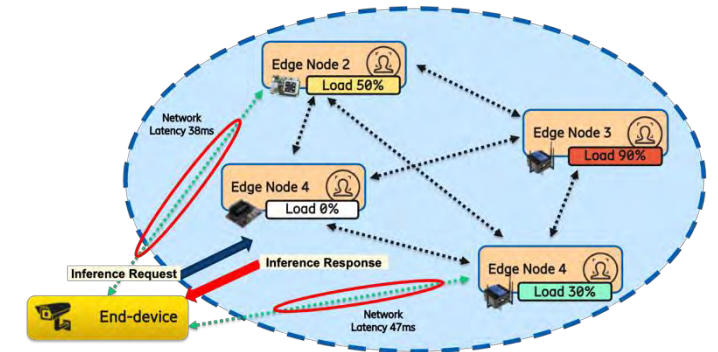
# Need For Newer Allocation Strategies

**Balancing Load and Network Latency in AI-Enabled Heterogeneous Edge Networks:** Proposes allocation based on device-specific CPU and AI accelerator performance, with a Performance Profiler assigning weights to reflect each node's capabilities for task distribution.

$$W_{\text{combined}} = \alpha \times W_{\text{cpu}} + \beta \times W_{\text{ai}} + \gamma \times W_{\text{nl}}$$



**Dynamic Adaptation to Workload:** Framework allocates tasks by assessing workload features and dynamically updates node weights to optimize task execution in response to changing system demands and resource availability.

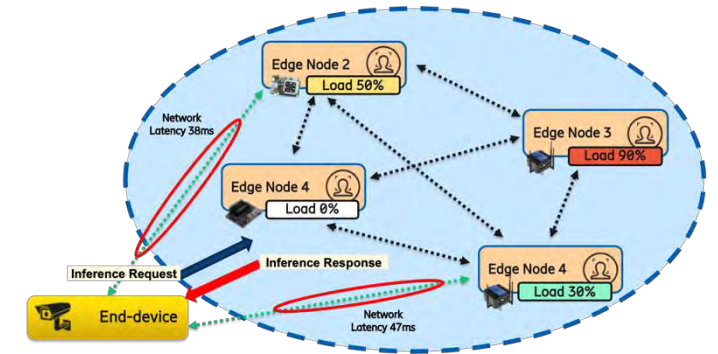


# Need For Newer Allocation Strategies

**Balancing Load and Network Latency in AI-Enabled Heterogeneous Edge Networks:** Proposes allocation based on device-specific CPU and AI accelerator performance, with a Performance Profiler assigning weights to reflect each node's capabilities for task distribution.

$$W_{\text{combined}} = \alpha \times W_{\text{cpu}} + \beta \times W_{\text{ai}} + \gamma \times W_{\text{nl}}$$

**FL angle?** What if  $\alpha$ ,  $\beta$ ,  $\gamma$  were learned in a federated way across nodes to reflect evolving conditions, while preserving privacy?



**Dynamic Adaptation to Workload:** Framework allocates tasks by assessing workload features and dynamically updates node weights to optimize task execution in response to changing system demands and resource availability.

# References

## Sources:

- Morabito R., and Chiang M., 2021, July. “Discover, Provision, and Orchestration of Machine Learning Inference Services in Heterogeneous Edge: A Demonstration”. In 2021 41st IEEE International Conference on Distributed Computing Systems (ICDCS 2021). IEEE. (**Best Demo Award**)
- Morabito R., Tatipamula M., Tarkoma S., and Chiang M., 2023. “Edge AI Inference in Heterogeneous Constrained Computing: Feasibility and Opportunities”. In 2023 IEEE International Workshop on Computer-Aided Modeling Analysis and Design of Communication Links and Networks (IEEE CAMAD).
- Morabito, R. and Chiang, M., 2024. Exploring Edge AI Inference in Heterogeneous Environments: Requirements, Challenges, and Solutions. In IoT Edge Intelligence (pp. 37-66). Cham: Springer Nature Switzerland.


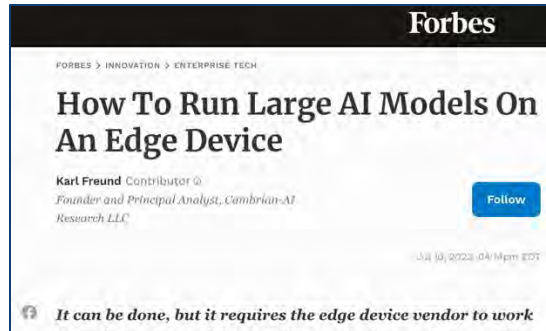
# **From the Edge to the Cloud: Exploring AI Inference Across the Computing Continuum**

(yes, including Generative AI)

**The Generative AI Wave: Is There Any Opportunity for the Edge?**

# The Generative AI Wave

## A Great Opportunity for the Edge



### MobileLLM: Optimizing Sub-billion Parameter Language Models for On-Device Use Cases

Zechun Liu, Changsheng Zhao, Forrest Iandola, Chen Lai, Yuandong Tian, Igor Fedorov, Yunyang Xiong, Ernie Chang, Yangyang Shi, Raghuraman Krishnamoorthi, Liangzhen Lai, Vikas Chandra


Sources:

<https://www.forbes.com/sites/karlfreund/2023/07/10/how-to-run-large-ai-models-on-an-edge-device/>  
<https://arxiv.org/pdf/2307.02779>  
<https://thenewstack.io/the-rise-of-small-language-models/>  
<https://arstechnica.com/information-technology/2024/04/apple-releases-eight-small-ai-language-models-aimed-at-on-device-use/>  
[https://www.youtube.com/watch?v=1DeQq1ZuG9o&t=155s&ab\\_channel=tinyMLFoundation](https://www.youtube.com/watch?v=1DeQq1ZuG9o&t=155s&ab_channel=tinyMLFoundation)  
<https://ieeexplore.ieee.org/document/10384606/>



# The Generative AI Wave

## A Great Opportunity for the Edge




Source:


<https://infohub.delltechnologies.com/en-us/p/investigating-the-memory-access-bottlenecks-of-running-llms/>

# The Generative AI Wave

## A Great Opportunity for the Edge



## Small Language Models (SLMs) Explosion




Source:

<https://infohub.delltechnologies.com/en-us/p/investigating-the-memory-access-bottlenecks-of-running-llms/>

# Large Language Models Inference Cost

Costs of inference as an application scales from 1k daily active users (DAUs) to 20k DAUs


- Each user sends an average of 15 requests per day




(Source: <https://www.inferless.com/learn/unraveling-gpu-inference-costs-for-llms-openai-aws-and-inferless>)

# Current Activities In This Area

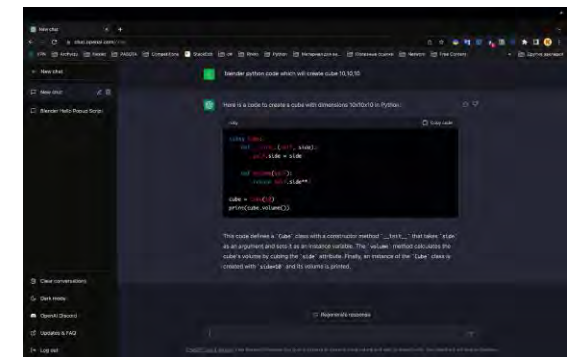
**Benchmarking SLMs in  
Constrained Devices**



**SLM/LLM Query Routing  
via Edge Collaboration**




**Streamlining TinyML  
Lifecycle with Large  
Language Models**



# Current Activities In This Area

**Benchmarking SLMs in  
Constrained Devices**



**SLM/LLM Query Routing  
via Edge Collaboration**



**Streamlining TinyML  
Lifecycle with Large  
Language Models**



# Benchmarking SLMs In Constrained Devices

## Small Language Models (SLMs) for resource constrained devices:

- Recent work shows the possibility to adopt LLMs at the constrained edge
- Still some way to go regarding the resource consumption when executed on MCU devices
- Potential to run on more capable but still lightweight edge devices (e.g., Single-Board Computers)




Source: <https://github.com/maxbbraun/llama4micro/>




# Benchmarking SLMs In Constrained Devices


**Current Focus:** developing a benchmark suite for evaluating the capabilities of LLMs on edge to **constrained edge devices**.




(a) RPi 5, Q4\_K\_M




(b) RPi 5, Q4\_0




(c) Orin GPU, Q4\_K\_M




(d) Orin GPU, Q4\_0




(a) Latency RPi 5, Q4\_K\_M




(b) Latency Orin CPU, Q4\_K\_M




(c) Latency Orin GPU, Q4\_K\_M



(d) Throughput RPi 5, Q4\_K\_M



(e) Throughput Orin CPU, Q4\_K\_M



(f) Throughput Orin GPU, Q4\_K\_M



# Read the Paper about this Benchmarking Work!

## **Sometimes Painful but Certainly Promising: Feasibility and Trade-offs of Language Model Inference at the Edge**

MAXIMILIAN ABSTREITER, University of Helsinki, Finland

SASU TARKOMA, University of Helsinki, Finland

ROBERTO MORABITO, EURECOM, France and University of Helsinki, Finland

The rapid rise of Language Models (LMs) has expanded the capabilities of natural language processing, powering applications from text generation to complex decision-making. While state-of-the-art LMs often boast hundreds of billions of parameters and are primarily deployed in data centers, recent trends show a growing focus on compact models—typically under 10 billion parameters—enabled by techniques such as quantization and other model compression techniques. This paper explores the feasibility and trade-offs of running LMs on edge devices, focusing on the challenges of limited memory, processing power, and energy consumption. We evaluate the performance of various LMs on edge devices and discuss the implications for future research and deployment.

<https://arxiv.org/pdf/2503.09114>





# Read the Paper about this Benchmarking Work!


## Sometimes Painful but Certainly Promising: Feasibility and Trade-offs of Language Model Inference at the Edge

- (C1) Addressing (RQ1-3), we benchmark 11 generative LMs on two widely used SBCs, analyzing key performance indicators such as memory usage, inference speed, and energy efficiency to quantify the feasibility of edge inference.
- (C2) Answering (RQ2), we evaluate the effect of quantization and model scaling on inference efficiency, resource utilization, and model performance, highlighting the trade-offs between accuracy and computational cost.
- (C3) Directly addressing (RQ3), we compare CPU-based inference against GPU acceleration, investigating their impact on execution speed, energy consumption, and practical deployment feasibility on edge devices.
- (C4) Responding to (RQ4), we analyze the impact of power modes, threading configurations, and system settings, while also evaluating micro-architectural metrics, such as cache misses and context switches, to uncover bottlenecks and efficiency gaps in edge-based LM execution.
- (C5) Informed by (RQ5), we assess the practical challenges of running LMs at the edge, including inference cost analysis, qualitative benchmarking of the models, and aspects such as usability related to real-world applicability, providing a broader perspective beyond raw performance metrics.




# Current Activities In This Area

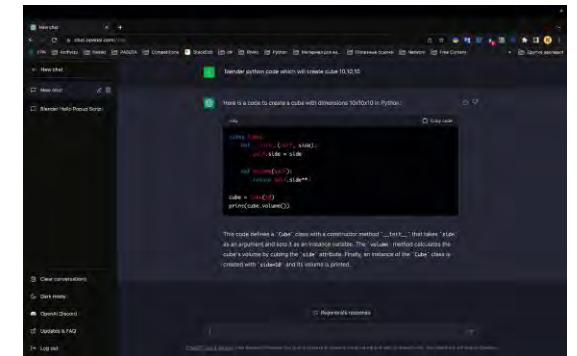
Benchmarking SLMs in  
Constrained Devices




SLM/LLM Query Routing  
via Edge Collaboration






Streamlining TinyML  
Lifecycle with Large  
Language Models



# LLMs **for** Edge AI: Any Possibility?



# Hardware – Key Differences

	Microprocessor	>	Microcontroller
<b>Platform</b>	 		
<b>Compute</b>	1GHz–4GHz	~10X	1MHz–400MHz
<b>Memory</b>	512MB–64GB	~10000X	2KB–512KB
<b>Storage</b>	64GB–4TB	~100000X	32KB–2MB
<b>Power</b>	30W–100W	~1000X	150μW–23.5mW

Source: Content on these slides is sourced from <https://github.com/edgeimpulse/courseware-embedded-machine-learning> and <https://github.com/tinyMLx/courseware/tree/master/edX>

# Portability Trade-offs

## MICROPROCESSOR


Universal Code Portability/Compatibility	✓
Cost (\$)	✗
Power Consumption (W)	✗
Engineering Effort	✗

## MICROCONTROLLER

Lower Code Portability	✗
Cost (\$)	✓
Power (W)	✓
Eng. Effort	✓


Source: Content on these slides is sourced from <https://github.com/edgeimpulse/courseware-embedded-machine-learning> and <https://github.com/tinyMLx/courseware/tree/master/edX>

# ML Software



```
import numpy as np
```

```
for x in range(10):  
    np.SaveTheWorld()
```

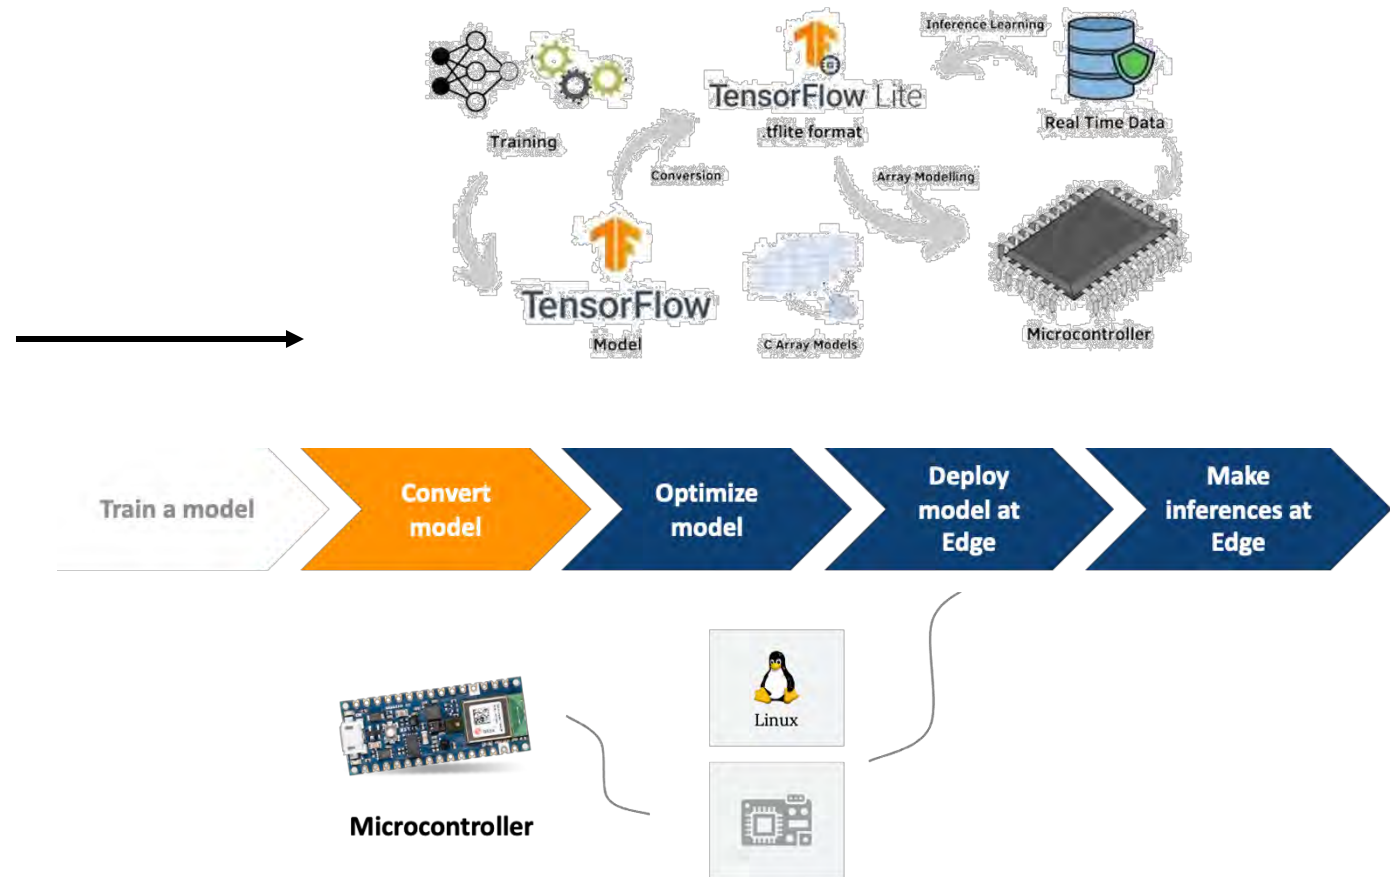



Able to execute the same code on different microprocessor hardware and architectures.

Source: Content on these slides is sourced from <https://github.com/edgeimpulse/courseware-embedded-machine-learning> and <https://github.com/tinyMLx/courseware/tree/master/edX>

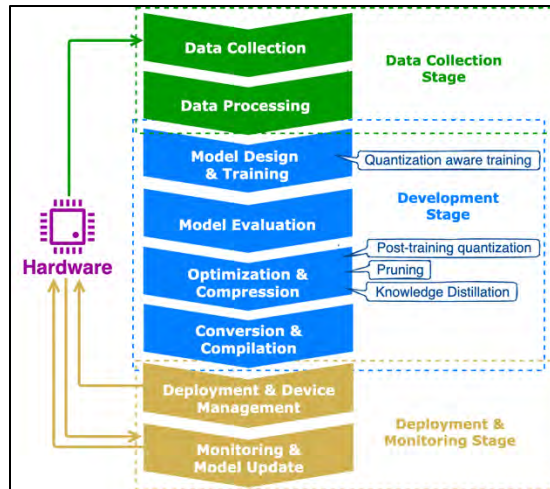
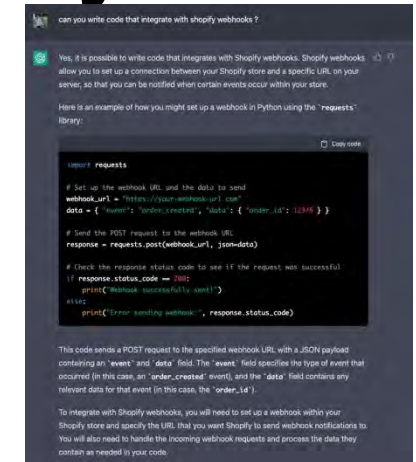
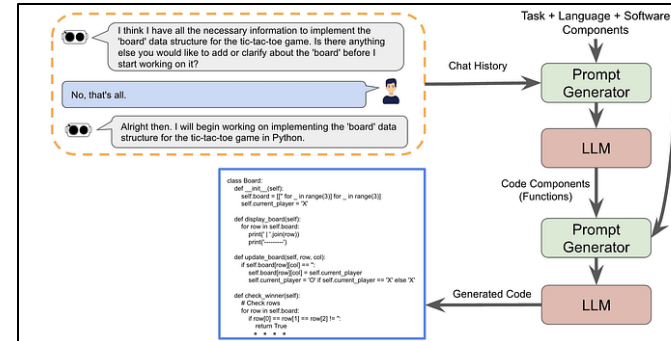


# ML Software For Constrained Devices




Source: Content on these slides is sourced from <https://github.com/edgeimpulse/courseware-embedded-machine-learning> and <https://github.com/tinyMLx/courseware/tree/master/edX>

# LLMs **for** Edge AI: Any Possibility?



**Edge AI Lifecycle**



## Code Generation Capabilities

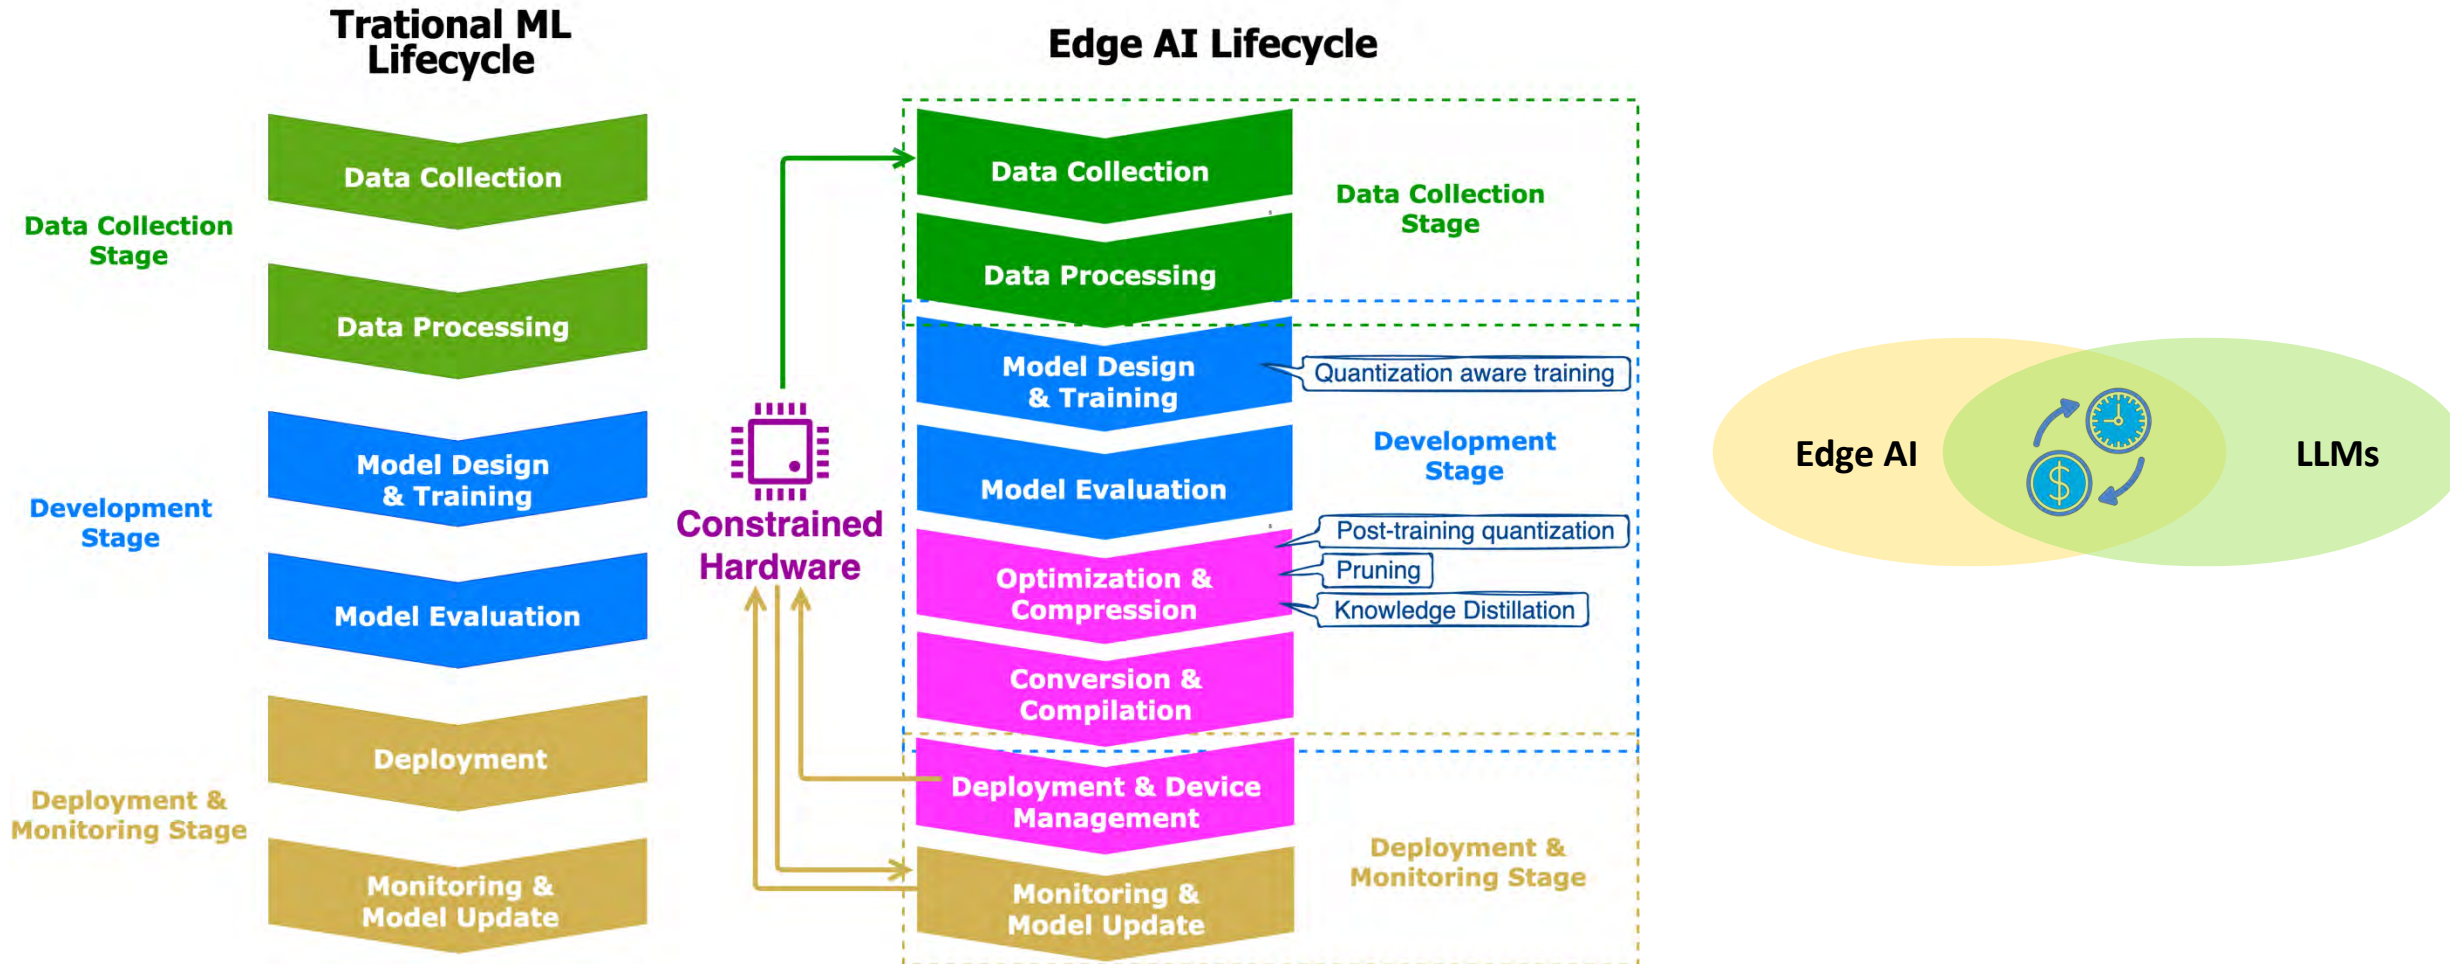
Model	Pricing	Pricing with Batch API*
gpt-4o	\$2.50 / 1M input tokens	\$1.25 / 1M input tokens
	\$1.25 / 1M cached** input tokens	
	\$10.00 / 1M output tokens	\$5.00 / 1M output tokens
gpt-4o-2024-08-06	\$2.50 / 1M input tokens	\$1.25 / 1M input tokens
	\$1.25 / 1M cached** input tokens	
	\$10.00 / 1M output tokens	\$5.00 / 1M output tokens
gpt-4o-2024-05-13	\$5.00 / 1M input tokens	\$2.50 / 1M input tokens
	\$15.00 / 1M output tokens	\$7.50 / 1M output tokens

## API Cost

Sources:

[https://medium.com/@saeedshamshiri\\_94060/looking-inside-gpt-synthesizer-and-the-idea-of-llm-based-code-generation-ff776b9e902f](https://medium.com/@saeedshamshiri_94060/looking-inside-gpt-synthesizer-and-the-idea-of-llm-based-code-generation-ff776b9e902f)  
<https://medium.com/@diegodursel/coding-with-chat-gpt-real-intelligence-b5e6ef129b4>  
<https://openai.com/api/pricing/>

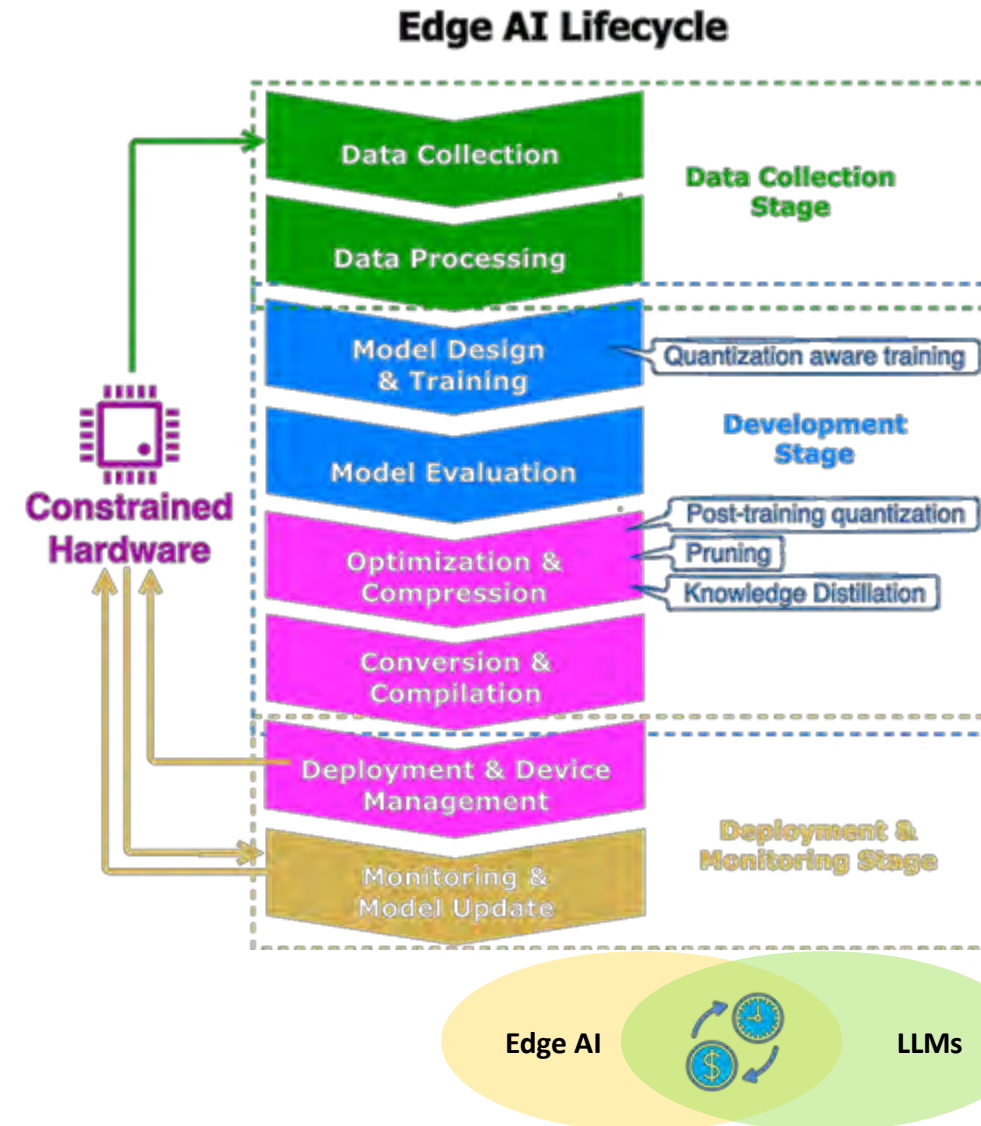
# LLMs for Edge AI Lifecycle Automation



# What Is Then This Work About?

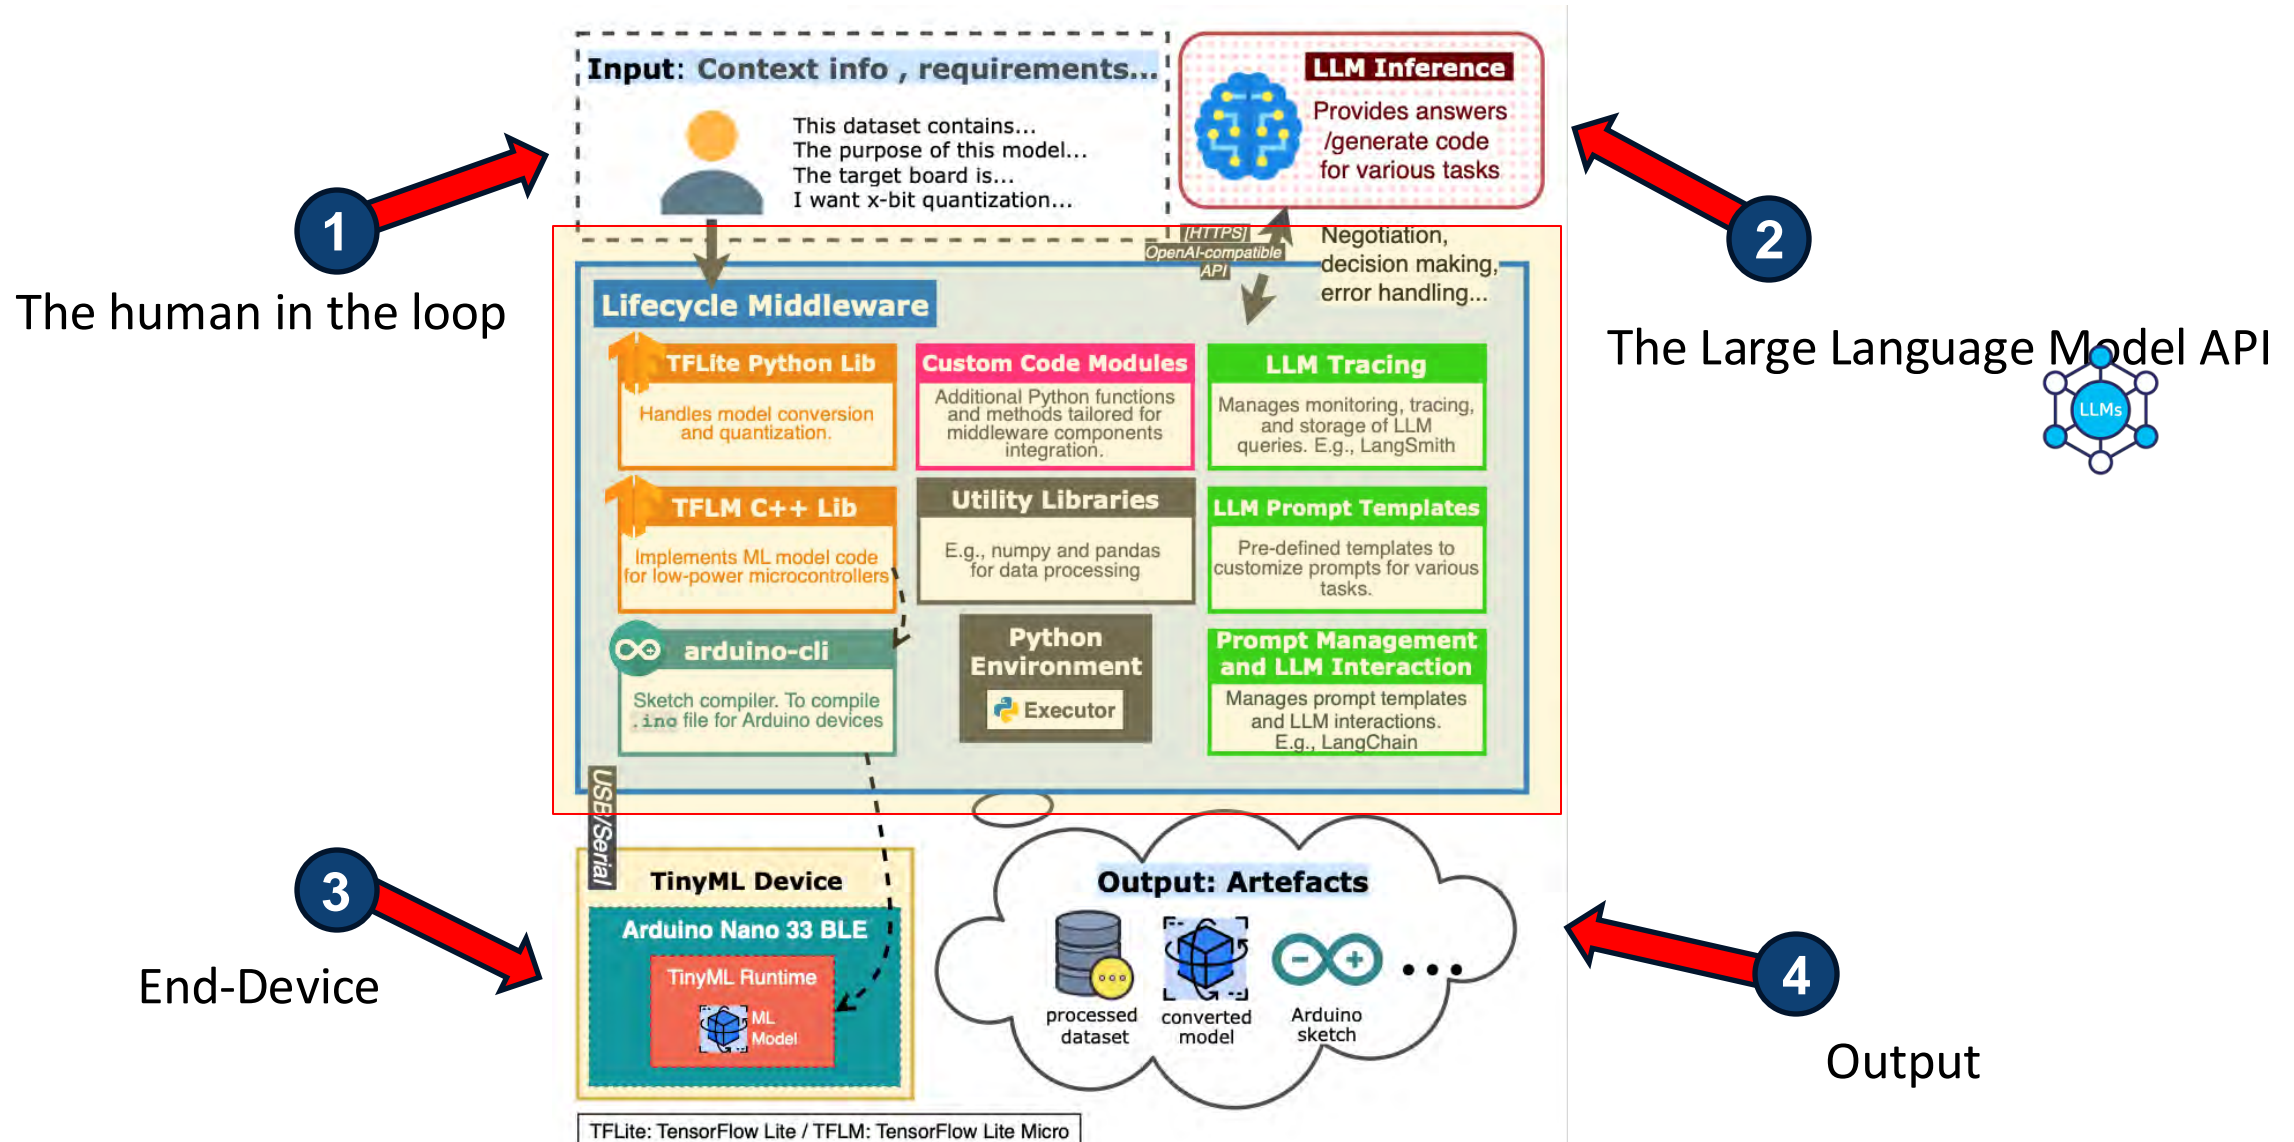
Three main questions we would like to answer:

1. What aspects of the **Edge AI lifecycle** can be processed and automated using LLMs?
2. How can **LLMs** be effectively **tailored** to optimize **Edge AI lifecycle stages**?
3. What are the **trade-offs, challenges**, and real-world considerations when integrating LLMs with *EdgeAIOps*?



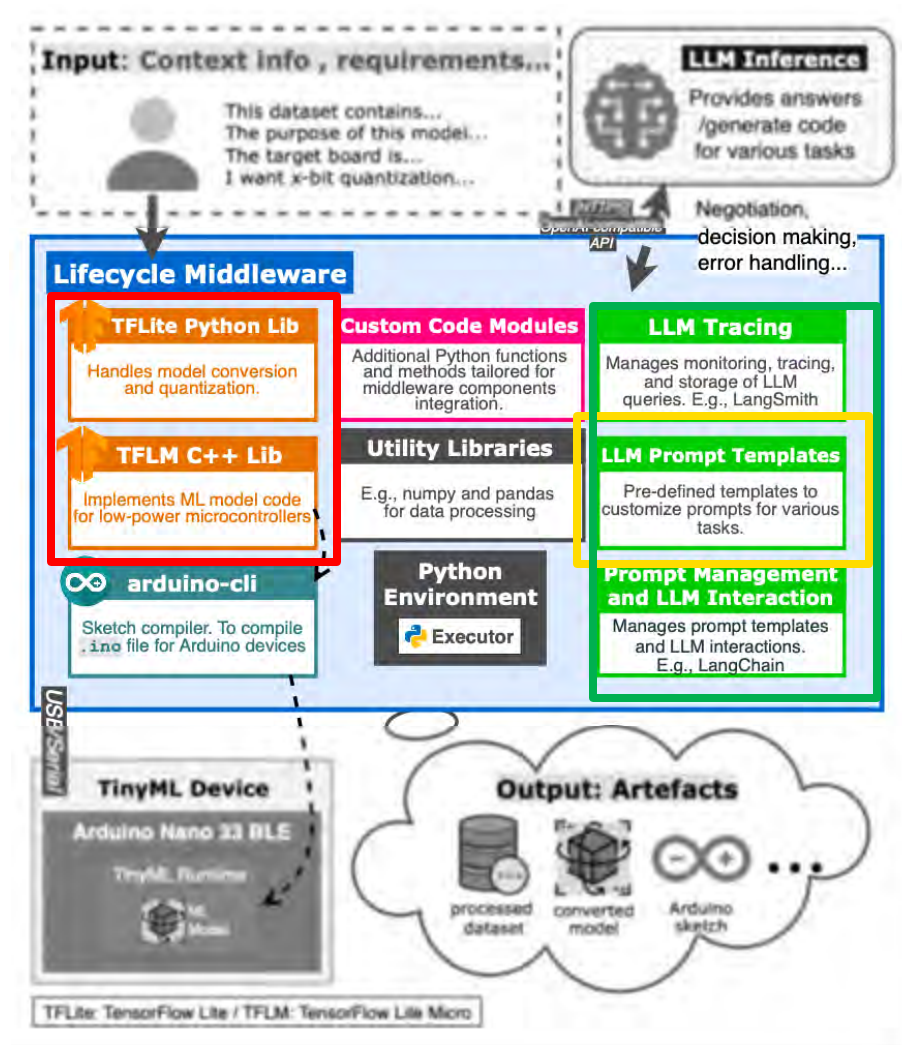


# The Overall View of Our Framework





# The Core of Our Framework



**LangChain** is a framework designed to build applications powered by LLMs that can connect to external data sources and perform dynamic decision-making based on those interactions.

**LangSmith** is a tool for debugging, testing, and monitoring LLM applications.

**LLMs Supporting Tools**

Predefined structure or format used to guide the LLMs' responses by embedding variables or placeholders within a fixed text – It helps ensuring consistent and targeted outputs for specific tasks.


**Prompt Templates**

```
>_ [green bar]
[blue bar] {{ [blue bar] }}
[blue bar]
```

# What About Prompts Templates?

Predefined structure or format used to guide the LLMs' responses by embedding variables or placeholders within a fixed text – It helps ensuring consistent and targeted outputs for specific tasks.

## Prompt Templates




## Prompt Engineering

(Effective communication & collaboration with AI)

Prompt Engineering is art and science of crafting inputs(prompts) to AI models to get the desired output

### Prompt Components



### Techniques

Zero-shot  
One-shot  
Few-shot  
Chain of Thought  
Self Consistency  
Generate Knowledge  
Automatic Prompt Engineering  
Active Prompt  
Directional Stimulus  
ReAct  
Multimodal CoT  
Graph Prompting

### Use Cases

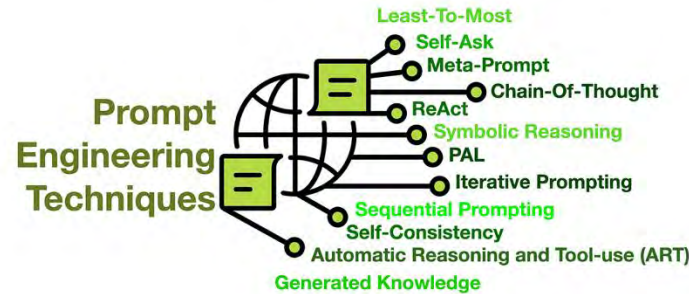
Text Summarization  
Question Answering  
Code Generation  
Role Playing  
Text Classification  
Reasoning  
Art Generation  
Grammar correction  
Bug finding  
Language Translation  
Idea Generation  
& many more

### Best Practices

Understand the model's capabilities and limitations  
Use clear and specific language  
Provide examples and feed  
Explain the context in as much detail as possible  
Experiment with different formats and styles  
Evaluate and refine

Source: <https://www.linkedin.com/pulse/importance-prompt-engineering-natural-language-c-cardoso-r->

## 12 Prompt Engineering Techniques



Source: <https://cobusgreyling.medium.com/12-prompt-engineering-techniques-644481c857aa>

www.cobusgreyling.com

## A Systematic Survey of Prompt Engineering in Large Language Models: Techniques and Applications

Pranab Sahoo<sup>1</sup>, Ayush Kumar Singh<sup>1</sup>, Sriparna Saha<sup>1</sup>, Vinija Jain<sup>2,3</sup>, Samrat Mondal<sup>1</sup> and Aman Chadha<sup>2,3</sup>

<sup>1</sup>Department of Computer Science And Engineering, Indian Institute of Technology Patna  
<sup>2</sup>Stanford University, <sup>3</sup>Amazon AI  
(pranab\_2021cs25, ayush\_2211ai27, sriparna, samrat@iitp.ac.in, hi@vinija.ai, hi@aman.ai)

### Abstract

Prompt engineering has emerged as an indispensable technique for extending the capabilities of large language models (LLMs) and vision-language models (VLMs). This approach leverages task-specific instructions, known as prompts, to enhance model efficacy without modifying the core model parameters. Rather than updating the model parameters, prompts allow seamless integration of pre-trained models into downstream tasks by eliciting desired model behaviors solely based on the given prompt.




Figure 1: Visual breakdown of prompt engineering components: LLMs trained on extensive data, instruction and context as pivotal elements shaping the prompt, and a user input interface.

Source: <https://arxiv.org/pdf/2402.07927>

## Forbes

FORBES > INNOVATION > AI

## The Best Prompt Engineering Techniques For Getting The Most Out Of Generative AI

Lance Eliot Contributor @

Dr. Lance B. Eliot is a world-renowned expert on Artificial Intelligence (AI) and Machine Learning...

Follow

May 9, 2024, 11:15am EDT

Source:

<https://www.forbes.com/sites/lanceeliot/2024/05/09/the-best-prompt-engineering-techniques-for-getting-the-most-out-of-generative-ai/>

# What About Prompts Templates?

Predefined structure or format used to guide the LLMs' responses by embedding variables or placeholders within a fixed text – It helps ensuring consistent and targeted outputs for specific tasks.

## Prompt Templates

```
full_pro_tem = """
### CONTEXT ###\n{context_prompt}
\n### OBJECTIVE ###\n{task_prompt}
"""

context_pro_tem = """You are an expert in Tiny Machine Learning (TinyML). \
highly skilled in the workflow, tools, techniques, \and best practices of TinyML operations. \
Your expertise extends to hardware, including microcontrollers. You will be asked questions \
regarding various tasks of TinyML, for example, data engineering, model designing, \
model evaluation, model conversion, deployment sketch developing, etc, of TinyML and may need \
to generate code to execute corresponding tasks, for example, data cleaning, model training code, etc."""
```

**Modular prompt templates** developed for different TinyML lifecycle stages include:

1. **Context**: Define LLM's role and expertise.
2. **Task-Specific Instructions**: Tailored for Edge AI tasks.
3. **Error Handling**: Help correct execution errors.
4. **Specification**: Provides a template for application specifications, such as hardware, sensors, software, and their configurations.
5. **Sketch Guideline Templates**: Provides code generation guidelines.



# What About Prompts Templates?

Predefined structured responses by embedding fixed text – It outputs for specific tasks

```
# Context template for TinyML expert role
context_pro_tem: str = (
    """You are an expert in Edge AI, highly skilled in the workflow, tools, \
    techniques, and best practices of Edge AI operations. Your expertise extends to hardware, including microcontrollers. \
    You will be asked questions regarding various tasks of Edge AI, for example, data engineering, model designing, \
    model evaluation, model conversion, deployment sketch developing, etc, of Edge AI and may need \
    to generate code to execute corresponding tasks, for example, data cleaning, model training code, etc."""
)

# Format template for specification filling responses
format_spec_filling_pro_tem: str = r"""
### RESPONSE FORMAT ###
- Output only one code block
- In the code block, only put the updated \"application_specifications\":{ }, remove everything outside of it.
- The response should be clear, accurate and strictly following the target_goal. Instead of assuming things, skip anything you are unsure about the detail.
"""

# Template for filling application specifications
# placeholders: board_fullname, dataset_summary, app_spec_pro_tem
# count_placeholders: 3
spec_filling_pro_tem: str = r"""
### OBJECTIVE ###
- **TASK**: Fill in the requested fields in application_specifications.
- **TASK INSTRUCTIONS**:
Read the value of \"board_fullname\" under \"hardware\" in \"application_specifications\", \
based on the board info, application description, and sensors that will be used in the application, \
fill in proper parameters into the placeholder fields decide_when_generating_code_based_on_given_board_and_application_description \
and decide_when_generating_code_based_on_given_data_sample_and_application_description. Keep \"guideline\" as it is originally. \
Make sure the libraries imported are all compatible with the board {board_fullname}, instead of using the library \"Qualcomm_TensorFlowLite.h\", \
directly use the library \"TensorFlowLite.h\"."

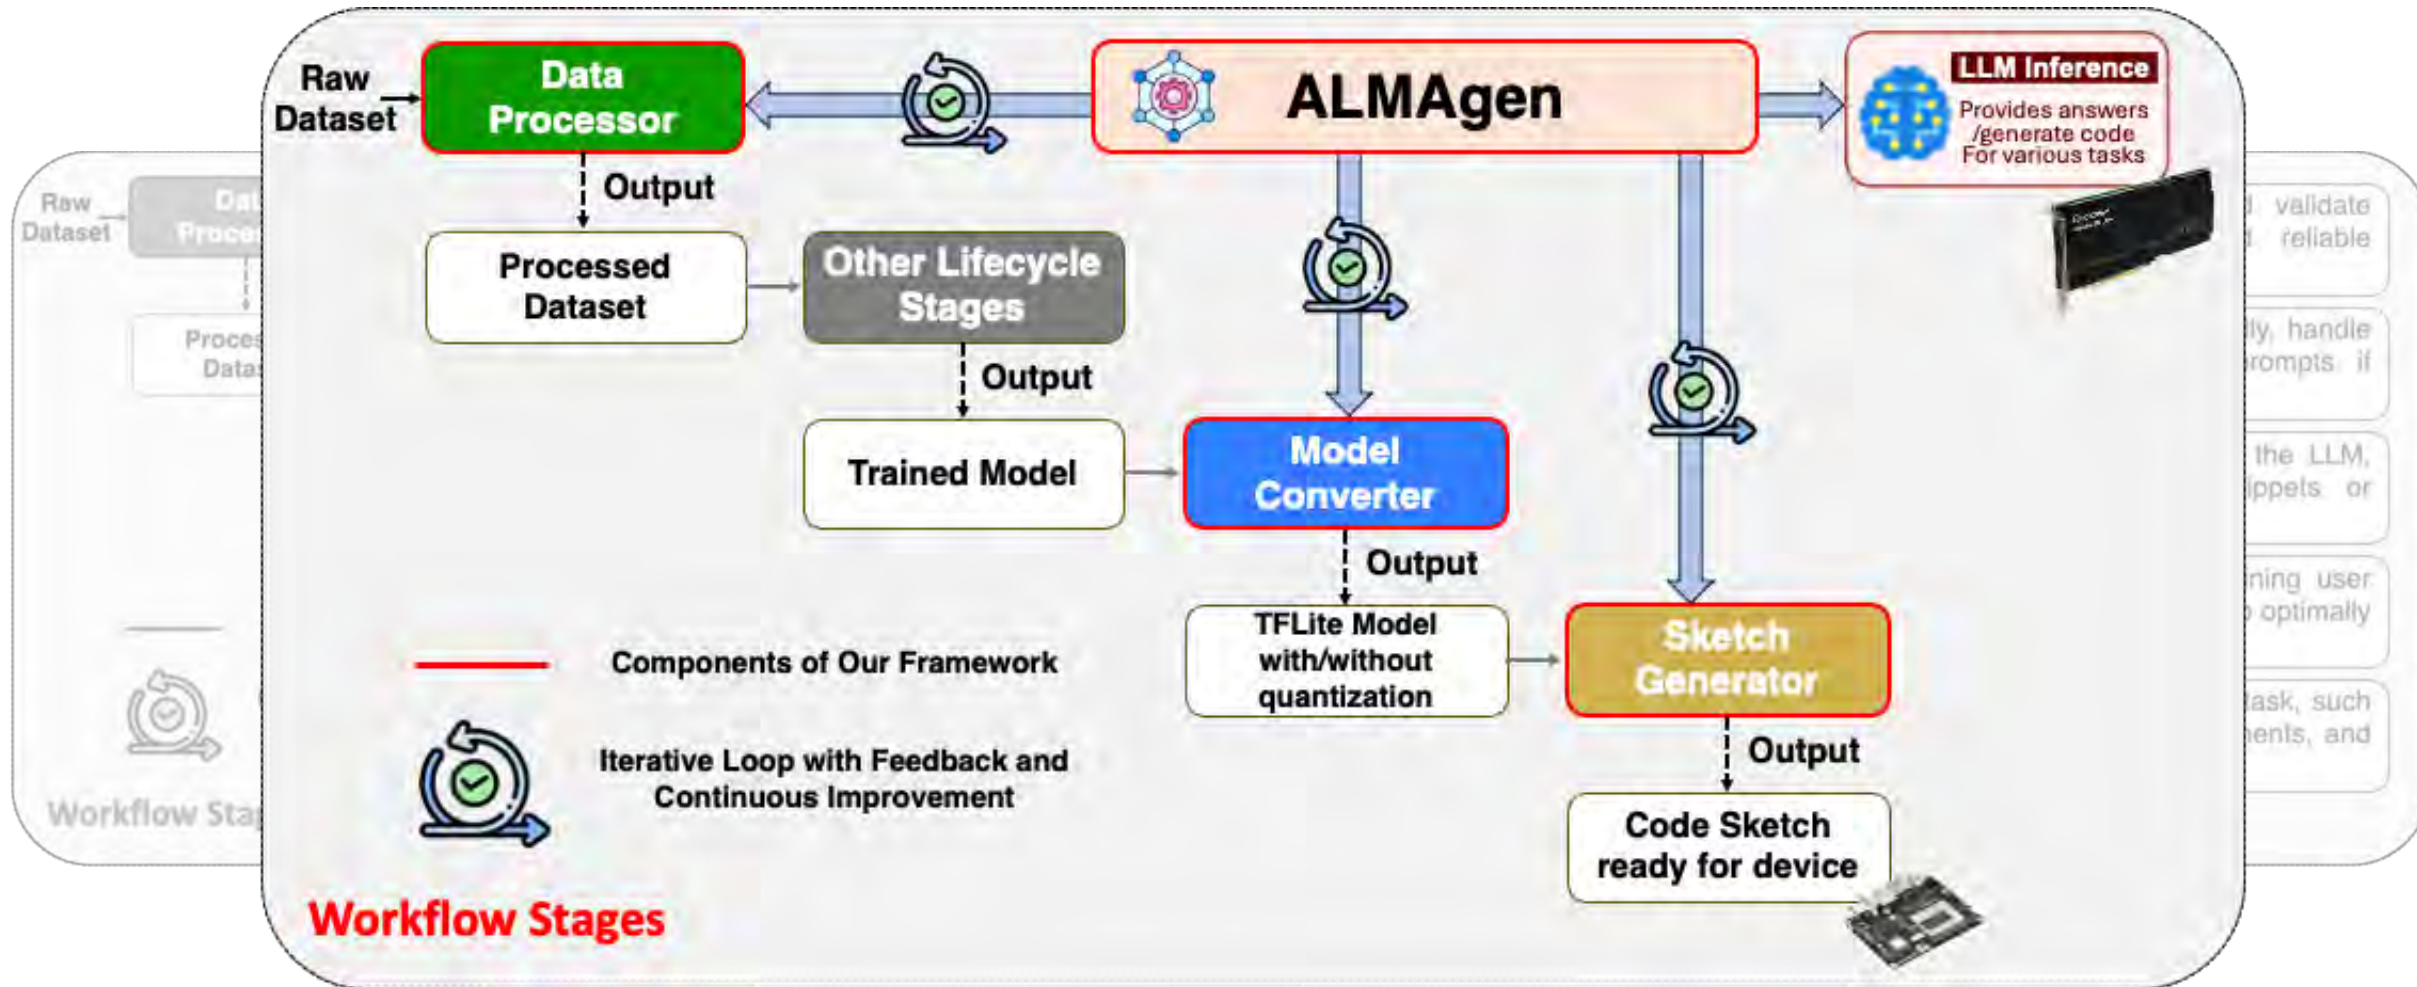
```

Context

Task

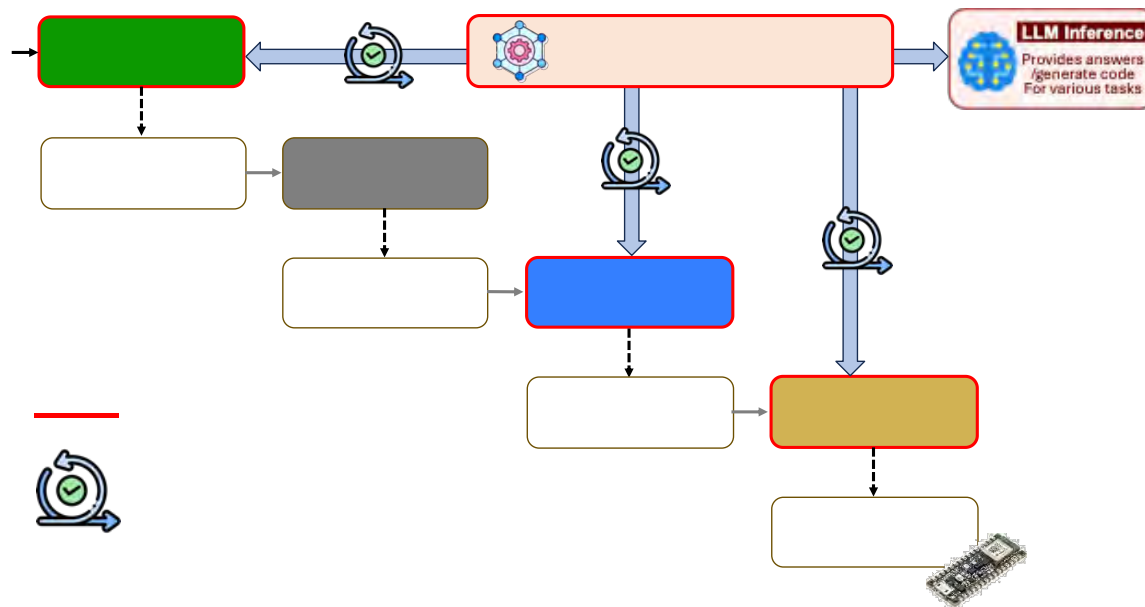
5. *Sketch Guideline Templates*: Provides code generation guidelines.

# Framework Workflow





# Test Case – CNN-based vision model



## Fruit identification using Arduino and TensorFlow

Posted by: ARDUINO TEAM — November 7th, 2019

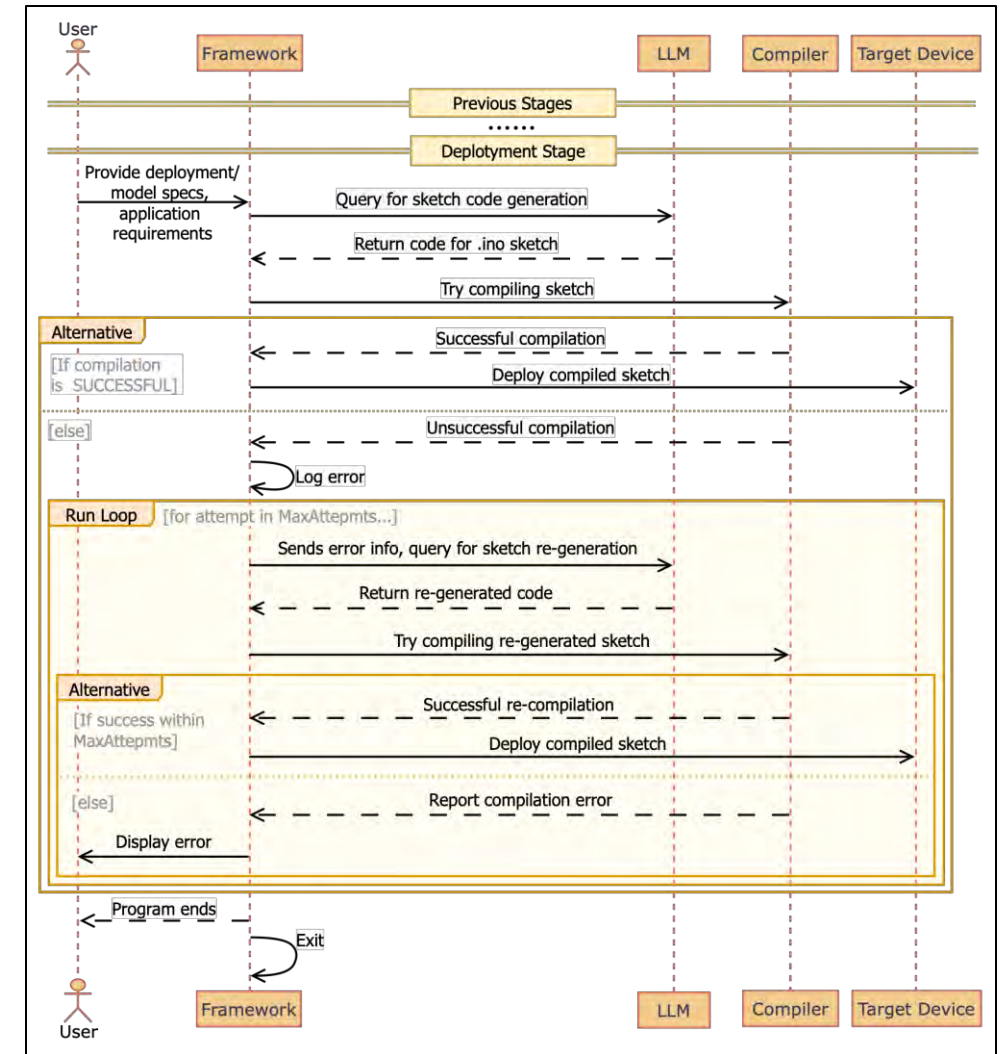
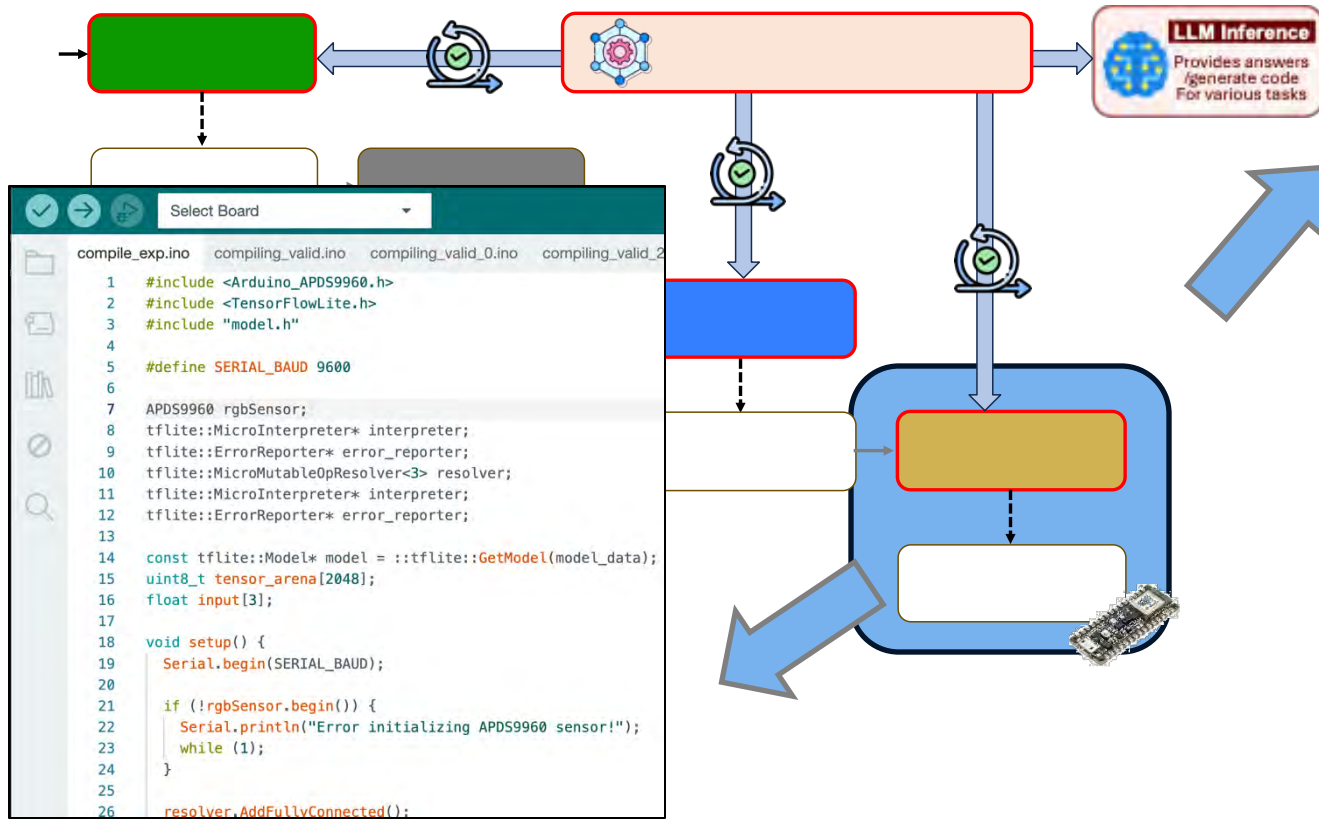
By Dominic Pajak and Sandeep Mistry

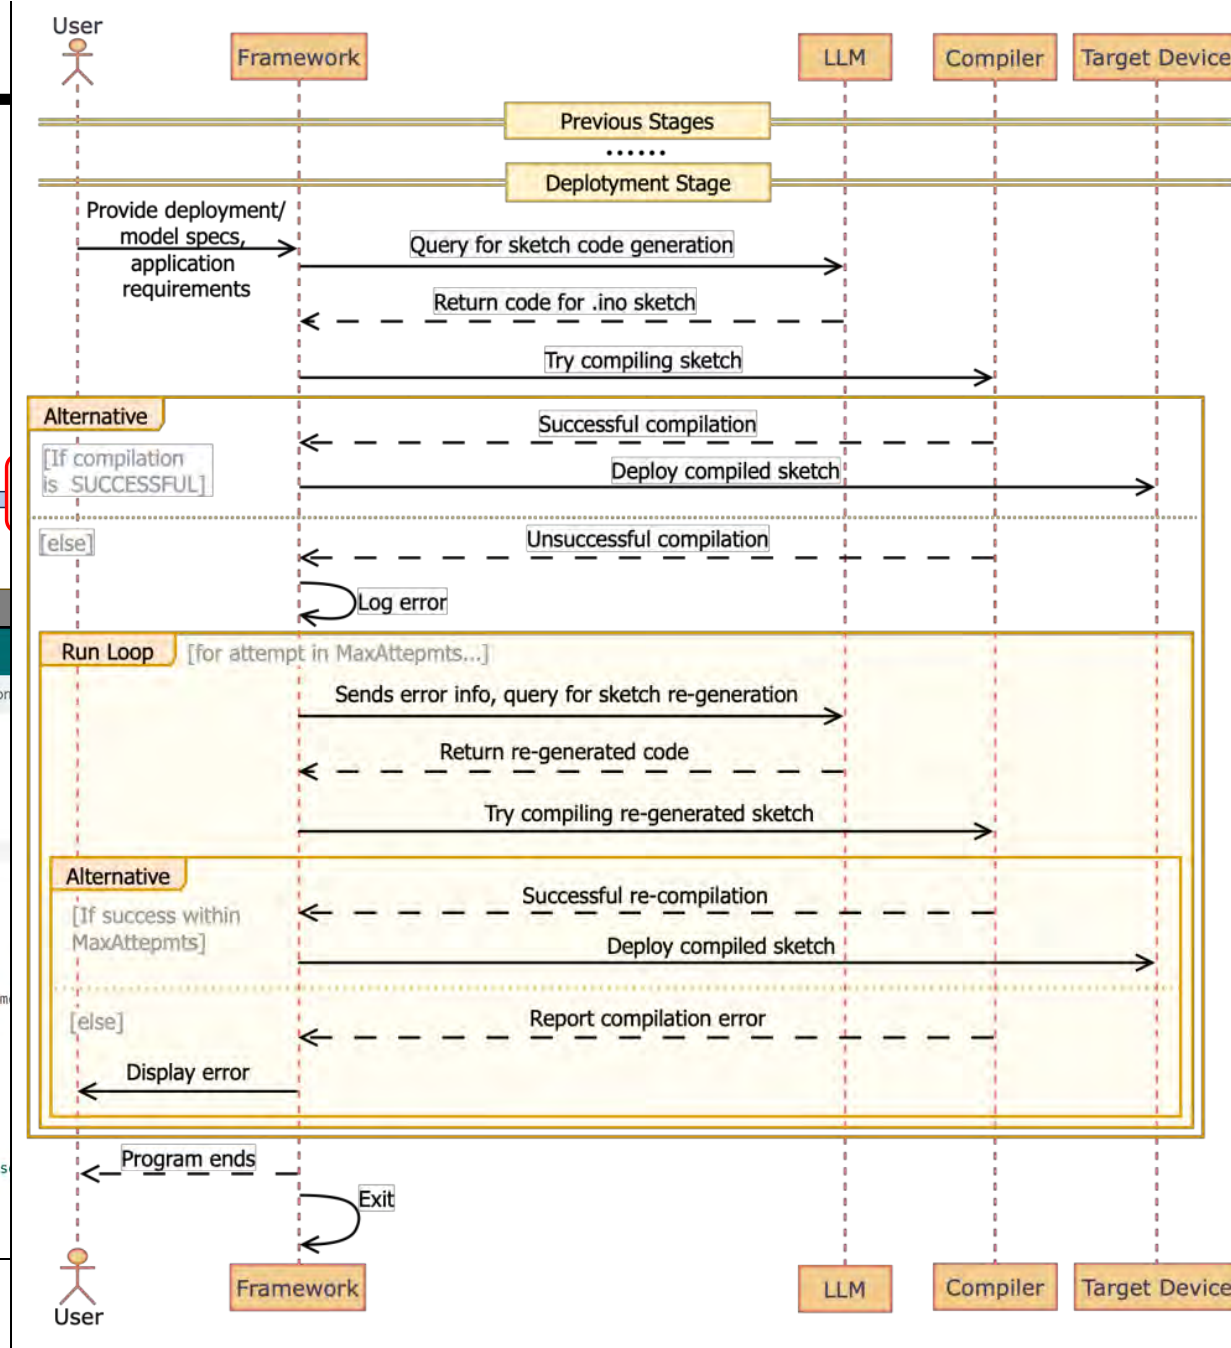
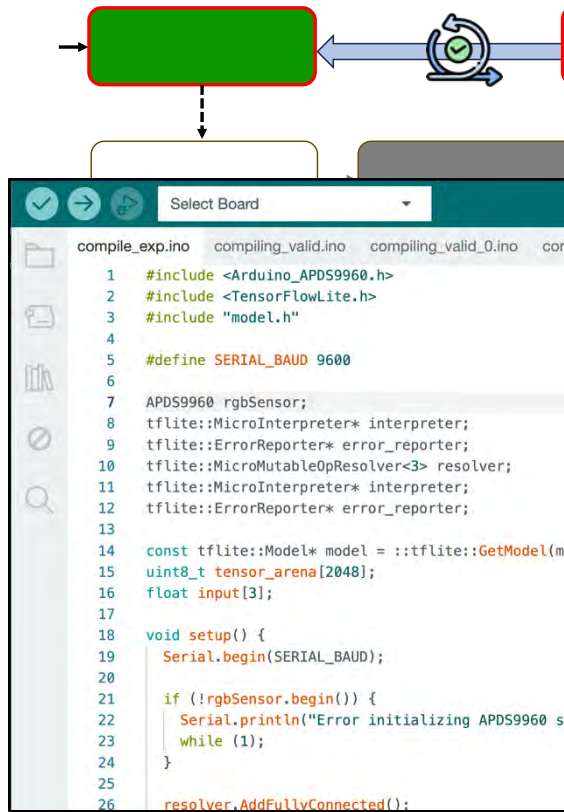


(Source: <https://blog.arduino.cc/2019/11/07/fruit-identification-using-arduino-and-tensorflow/>)

# Test Case – CNN-based vision model

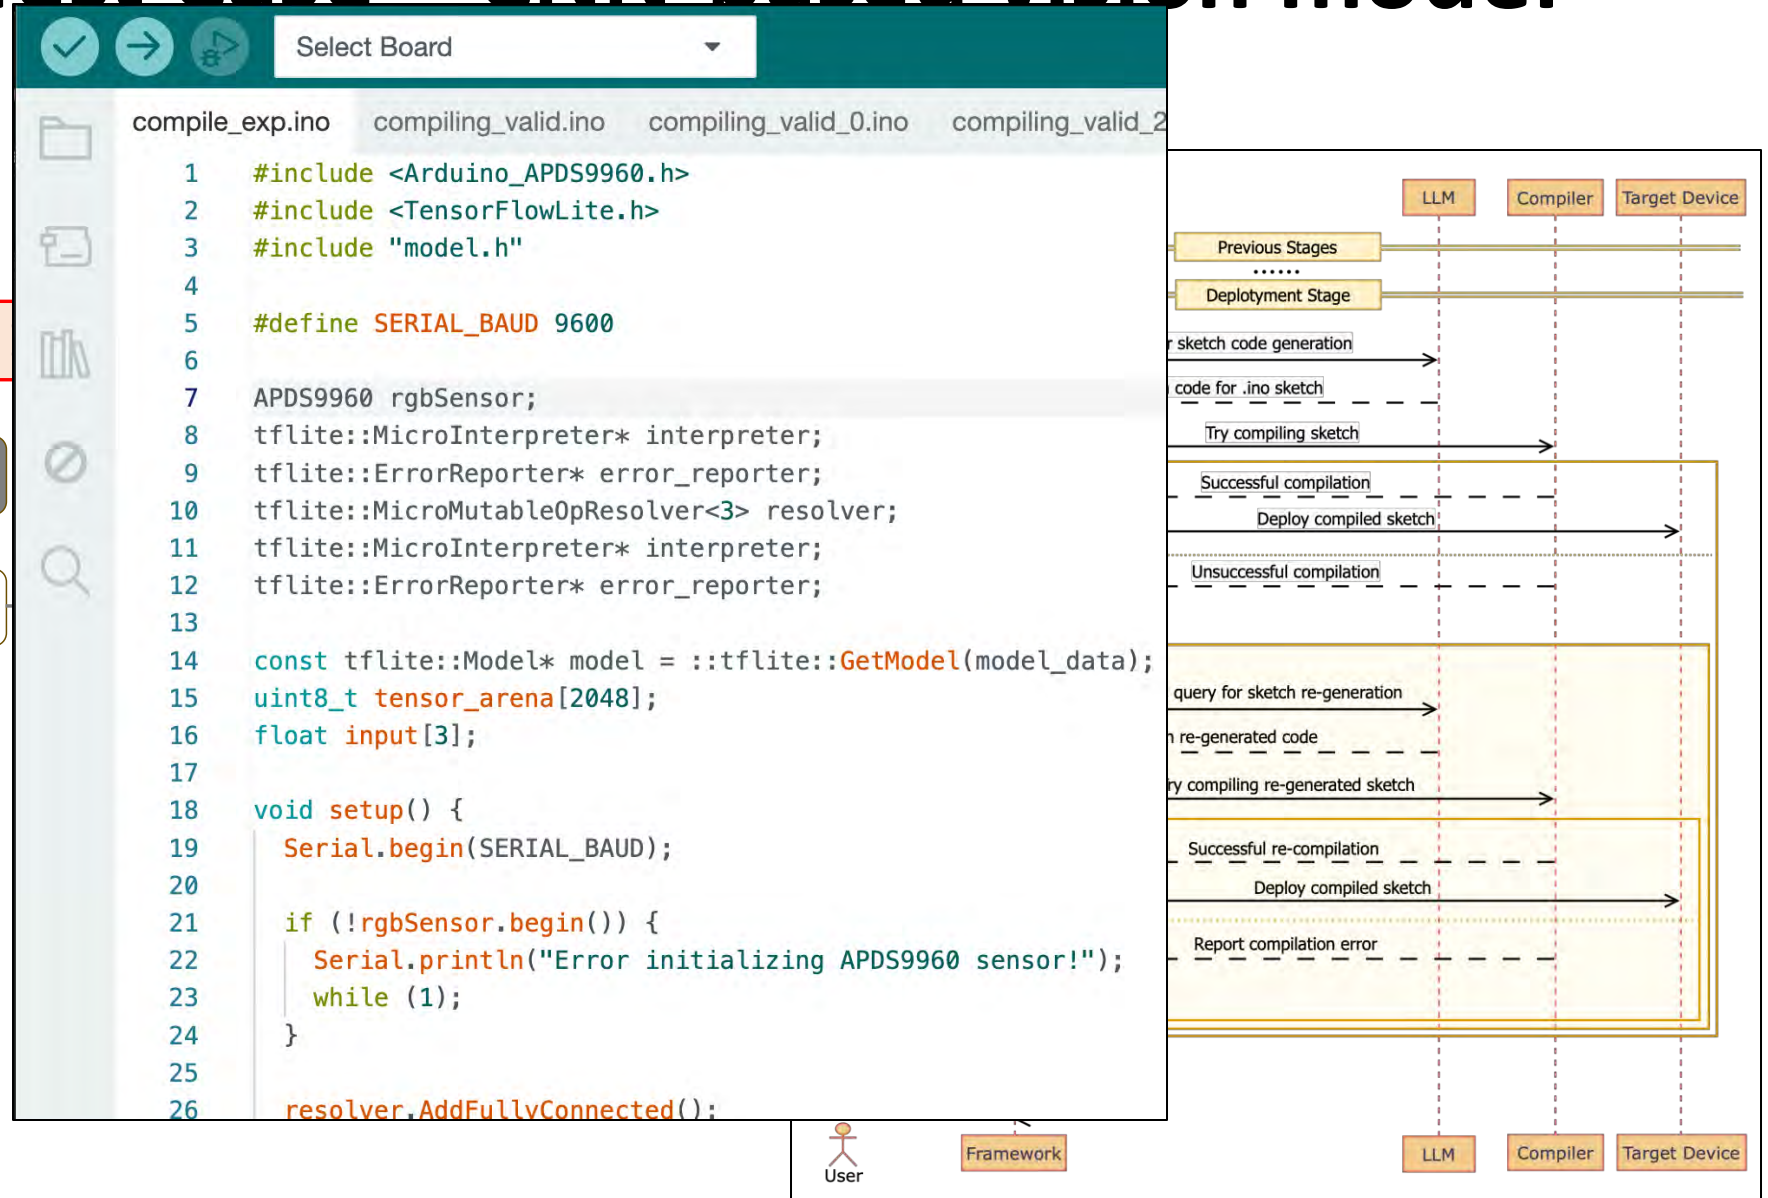
## Sketch Generation Lifecycle



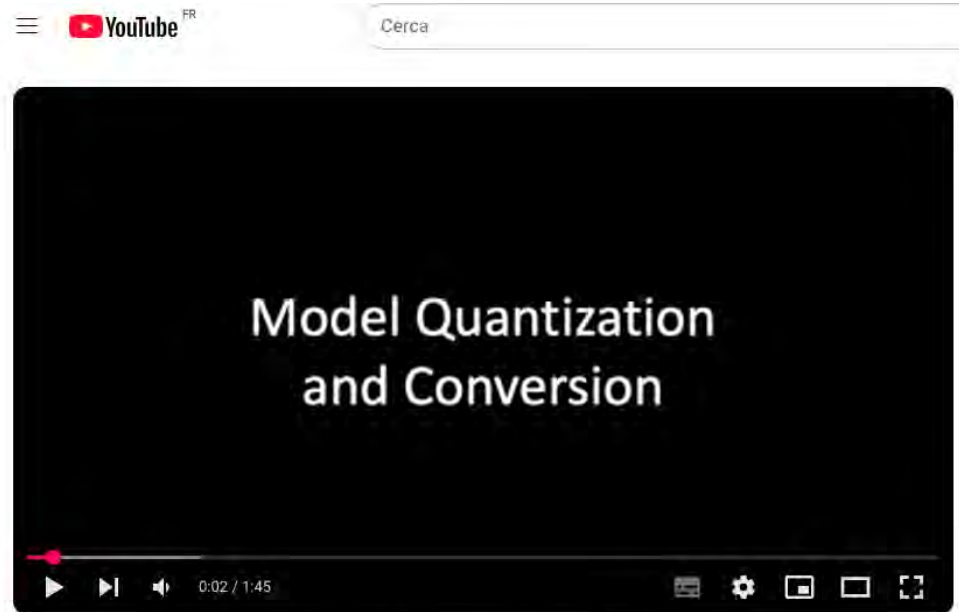


model





# Demos



**Automating Model Quantization and Conversion for TinyML with LLMs**

Non in elenco

**Roberto Mo...**  
2 iscritti

Analytics

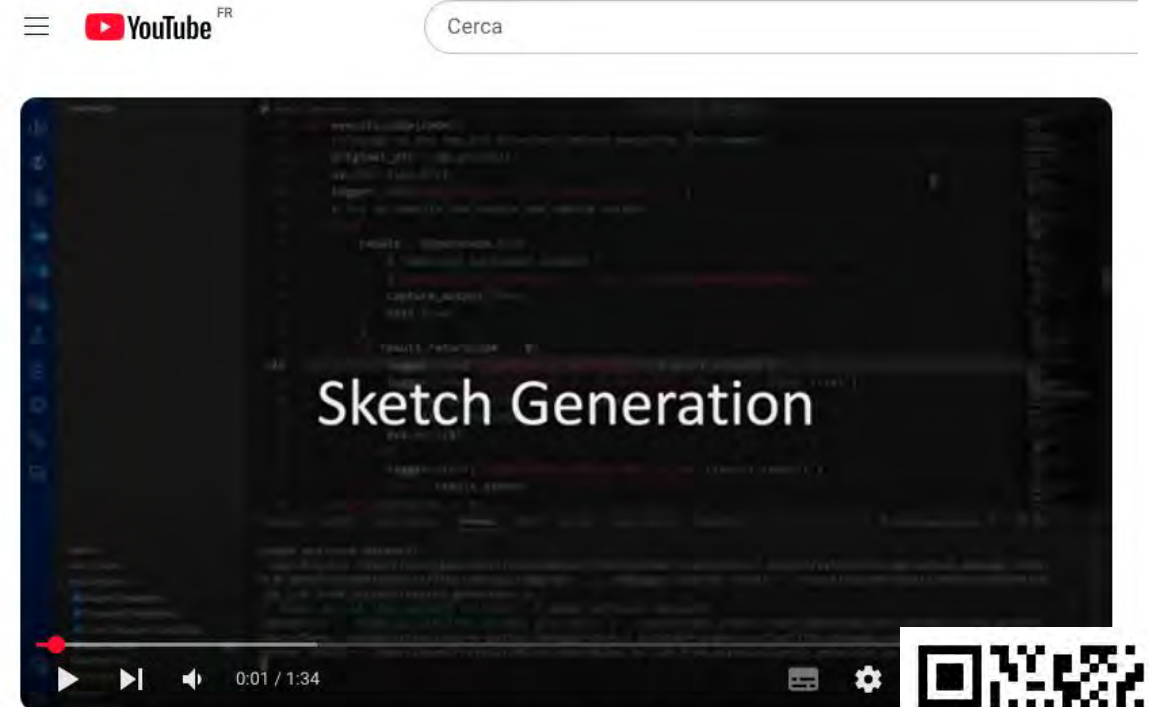
Modifica video

0

Condividi

...

[https://www.youtube.com/watch?v=KnJ5m78x\\_X8](https://www.youtube.com/watch?v=KnJ5m78x_X8)



**Automated Sketch Code Generation for TinyML on Arduino with LLMs**

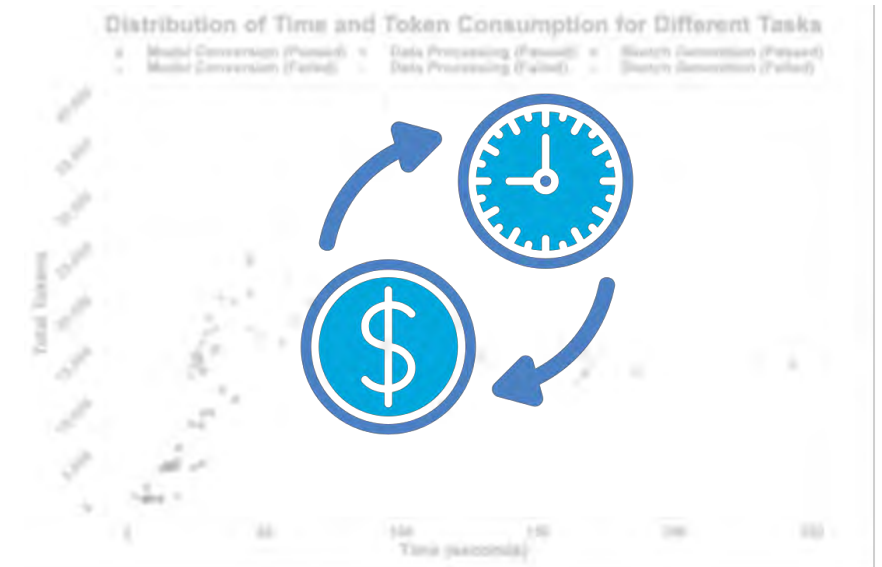

Non in elenco

<https://www.youtube.com/watch?v=Qjpsb5Wnnl8>





# What is The Cost of This?



*The data presented in this empirical evaluation is subject to change as OpenAI models are frequently updated, which may impact results over time.*

# What is The Cost of This?


## Tokens and Time Perspective

- A **token** is a unit of text (e.g., word, subword, or character) that the model processes.
- **Input tokens** are the tokens derived from the text provided to the model for analysis or generation.
- **Output tokens** are the tokens generated by the model in response to the input, forming the predicted or generated text.

*Both input and output tokens impact processing time and computational resources.*

# What is The Cost of This?

## Tokens and Time Perspective

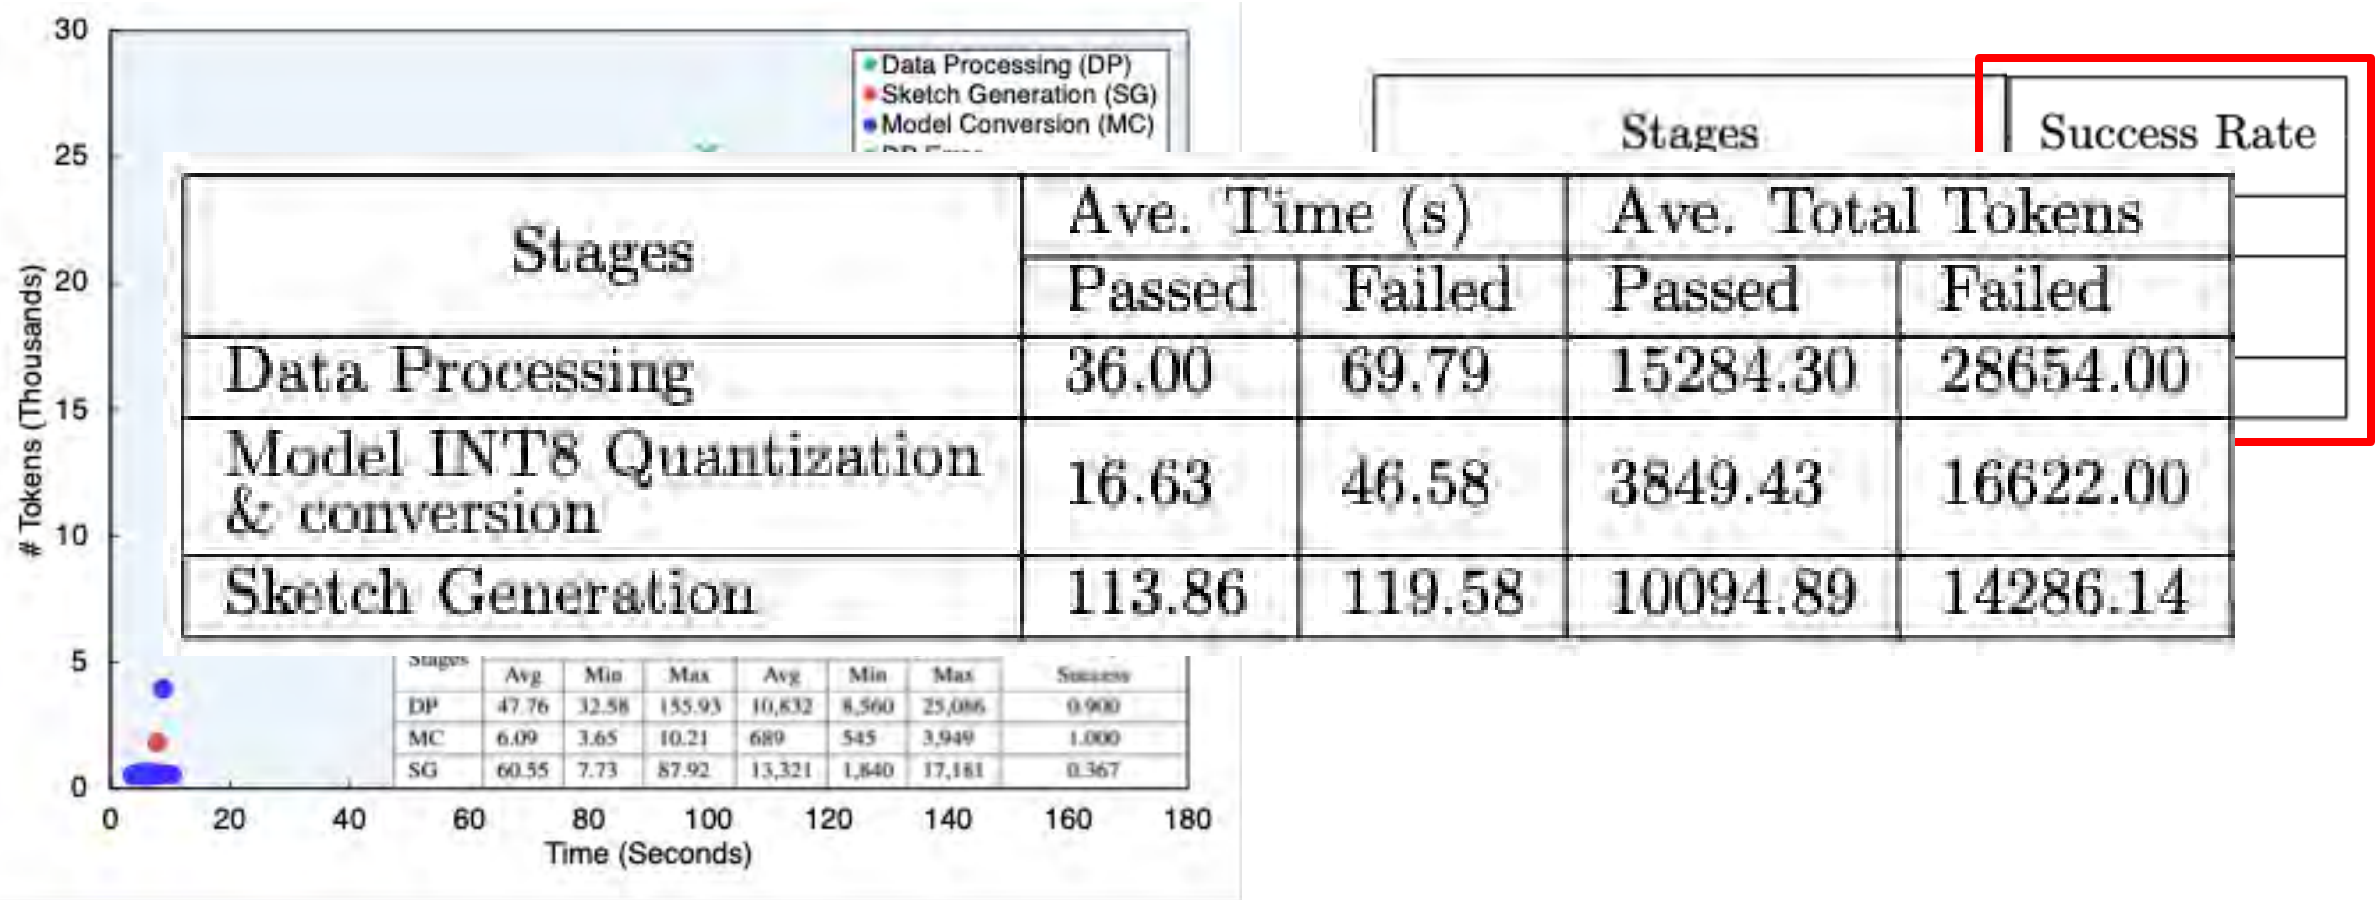


Stages	Success Rate
Data Processing	0.900
Model INT8 Quantization & conversion	0.933
Sketch Generation	0.300

Stages	Ave. Time (s)		Ave. Total Tokens	
	Passed	Failed	Passed	Failed
Data Processing	36.00	69.79	15284.30	28654.00
Model INT8 Quantization & conversion	16.63	46.58	3849.43	16622.00
Sketch Generation	113.86	119.58	10094.89	14286.14


# What is The Cost of This?

## Tokens and Time Perspective



# What is The Cost of This?

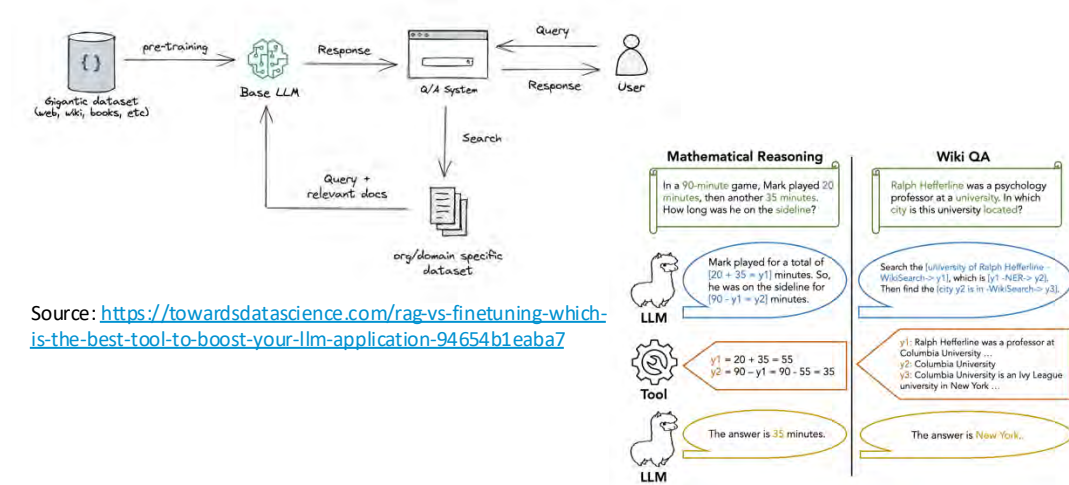
## Monetary Cost Perspective




# Reality, Illusion, or Opportunity?

Potential for expanded automation in the TinyML lifecycle.

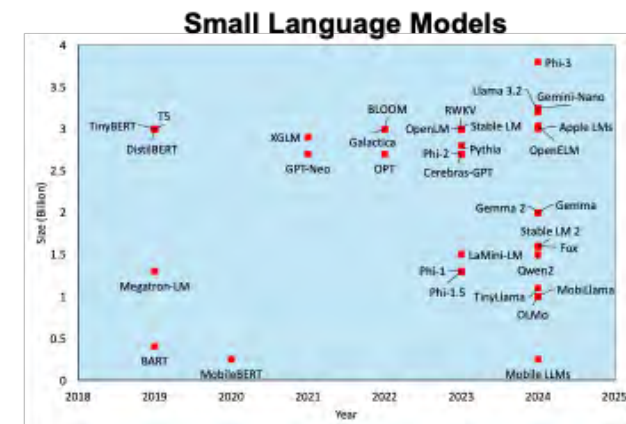
- **LLM fine-tuning** can improve code generation reliability we can enhance versatility.
- **Integrating LLMs with external tools** for enhanced reasoning can unlock additional level of reasoning (and so improvements).
- **Involving the end-device** may enable abstraction of device-specific info and real-time optimization for more efficient lifecycle management.



Source: <https://arxiv.org/pdf/2401.17464>



Source: <https://arxiv.org/pdf/2402.16844>





# Read the Paper about this Work!

## Consolidating TinyML Lifecycle with Large Language Models: Reality, Illusion, or Opportunity?

Guanghan Wu<sup>‡</sup>, Sasu Tarkoma<sup>‡</sup>, Roberto Morabito<sup>\*</sup>

<sup>\*</sup>Department of Communication Systems, EURECOM, France.

<sup>‡</sup>Department of Computer Science, University of Helsinki, Finland.

<https://arxiv.org/pdf/2501.12420>

To appear in IEEE IoT Magazine!




# Current Activities In This Area

Benchmarking SLMs in  
Constrained Devices



SLM/LLM Query Routing  
via Edge Collaboration



Streamlining TinyML  
Lifecycle with Large  
Language Models




# Towards An Efficient LLM Query Routing Via Edge Collaboration

**Problem:** Mobile or edge devices cannot support large LLMs due to their resource limitations, while cloud-based LLMs are expensive and raise privacy concerns.

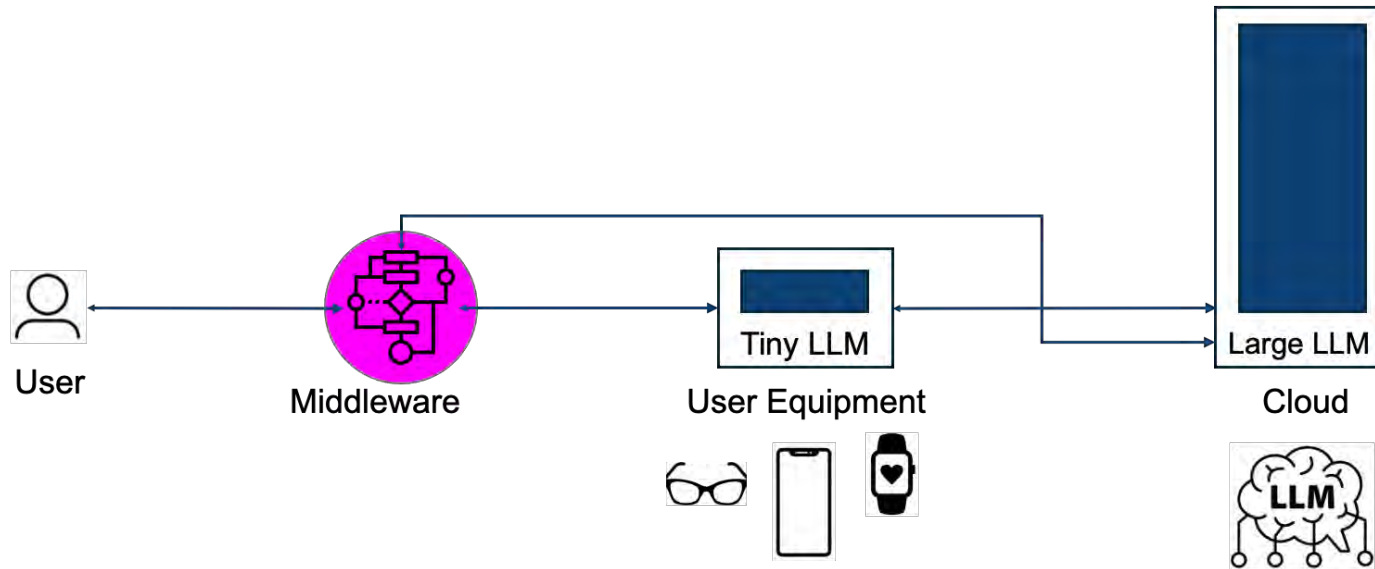
**Research Question:** How can we improve user query responses by collaboratively utilizing smaller local LMs and larger cloud-based LLMs?

## Challenge

- How do we determine which queries should be processed locally and which should be sent to the cloud?
- How can we balance the trade-offs between cost, performance, and privacy?



# Towards An Efficient LLM Query Routing Via Edge Collaboration

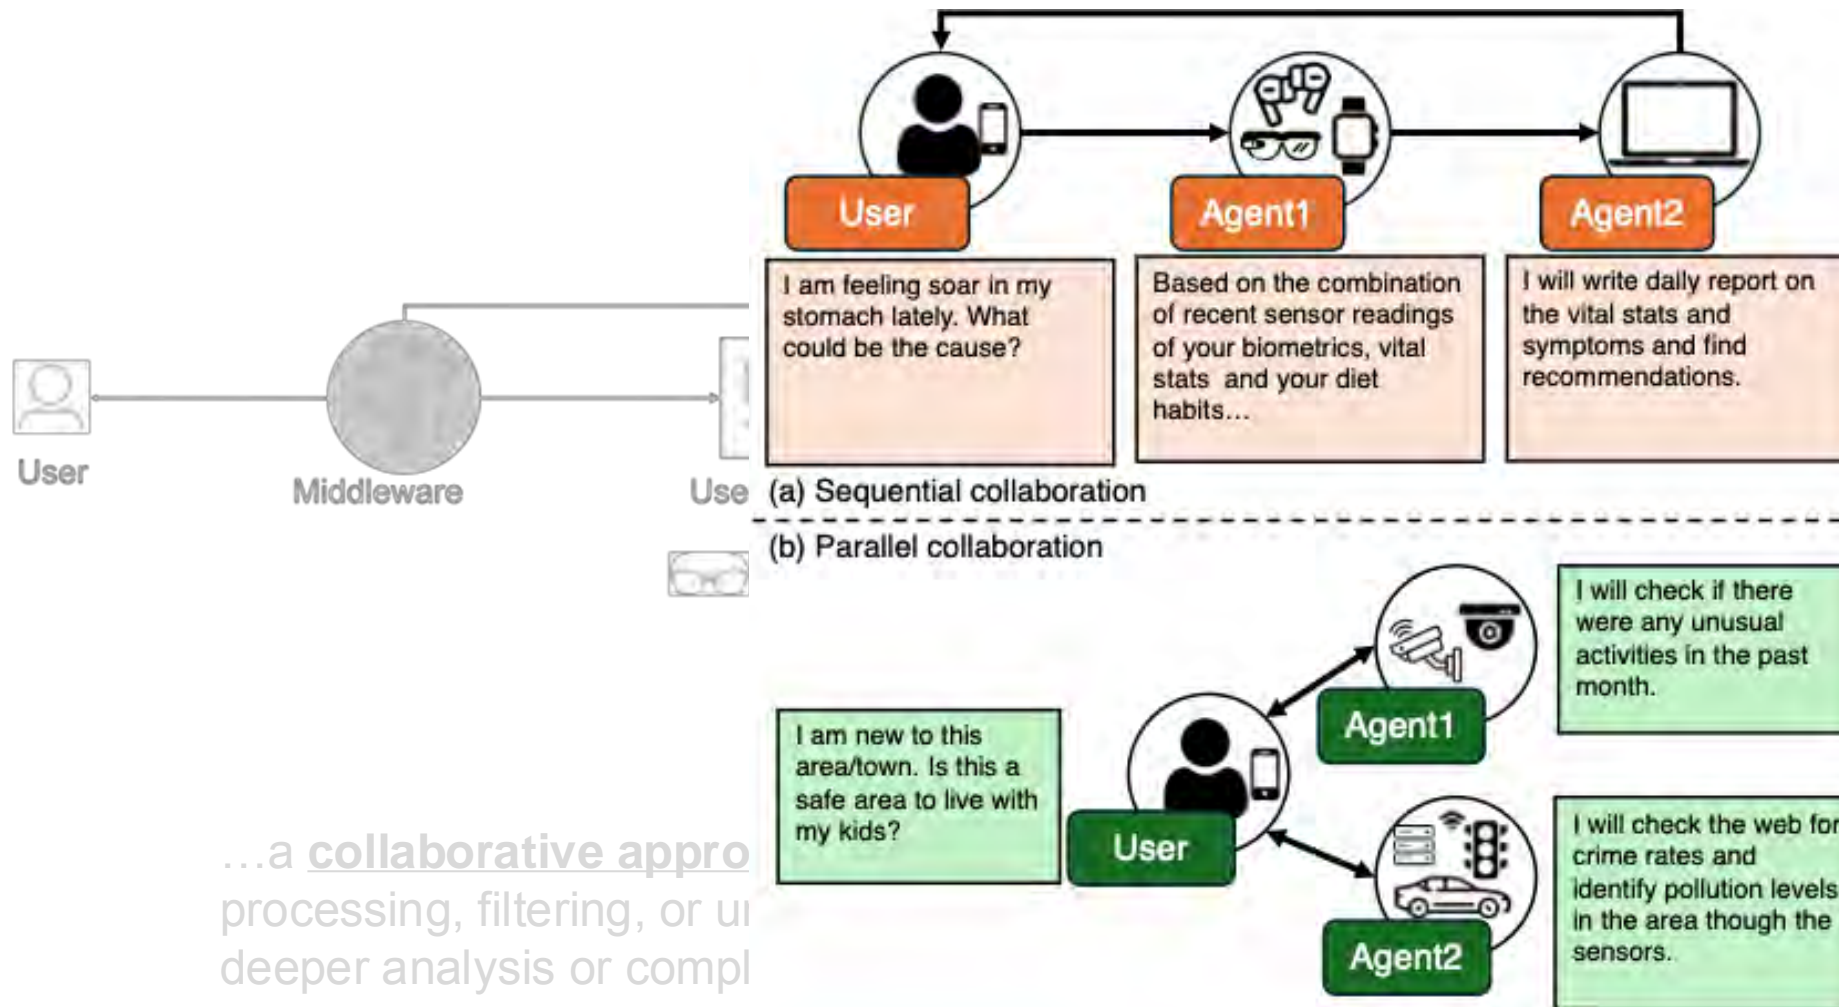


The key concept is ‘query split and routing.’

- If the query **complexity is high**, it is logical to route it to the back-end **cloud model**.
- If the query **complexity is low**, it makes sense to process it with the **front-end local model**.
- However, if the query falls into a **gray area** between local and back-end processing...

...a **collaborative approach** is ideal. In this case, the local model can handle preliminary processing, filtering, or understanding the context, while the cloud-based LLM provides deeper analysis or complex reasoning.

# Towards An Efficient LLM Query Routing Via Edge Collaboration



It is 'query split and

complexity is high, it is  
e it to the back-end

complexity is low, it  
to process it with the  
l model.


the query falls into a gray  
local and back-end

preliminary  
provides

...a collaborative approach  
processing, filtering, or u  
deeper analysis or compl



# Towards An Efficient LLM Query Routing Via Edge Collaboration



# Federated Learning Directions

## Federated Prompt Fine-Tuning

*What if prompt templates or instruction tuning for LLM-based Edge AI workflows were collaboratively adapted using FL?*

- Especially relevant when privacy prevents uploading prompt examples to a central server.

## FL for Model Routing Policies

*Could routing policies for SLM/LLM (i.e., when to run locally or send to the cloud) be learned in a federated fashion, adapting to each device's workload and usage patterns?*


- Leverages FL to train routing classifiers without leaking device usage data.

## FL for Lifecycle Feedback Loops

*Could lifecycle automation (e.g., code sketch generation) benefit from federated feedback loops, where each device refines LLM usage strategies based on local success/failure logs?*

- Could be posed as a federated reinforcement learning or federated policy optimization challenge.

### SLM/LLM Query Routing via Edge Collaboration



### Streamlining TinyML Lifecycle with Large Language Models



# From the Edge to the Cloud: Exploring AI Inference Across the Computing Continuum

(yes, including Generative AI)

Roberto Morabito  
Assistant Professor @ EURECOM  
<https://www.linkedin.com/in/robertomorabito>

Kudos to Maximilian Abstreiter, Guanghan Wu (University of Helsinki), and SiYoung Jang (Nokia Bell Labs)

# About EURECOM

## Student Activities

EURECOM has a very active Student's association organises many events all year long to fully enjoy your stay in France.



**Weekend of Integration**



**International Meal**



**International Study trip**



**Ski  
Weekend**



**Many sport activities**



**Graduation and  
Alumni Events**